Privacy and Polarization: An Inference-Based Framework

Tommaso Bondi,[a] Omid Rafieian,[a] Yunfei (Jesse) Yao[b],*

[a] Marketing, Cornell Tech and SC Johnson College of Business, Cornell University, New York, New York 10044; [b] The Chinese University of Hong Kong, Shatin, Hong Kong
*Corresponding author
Contact: tbondi@cornell.edu, https://orcid.org/0000-0003-4572-0709 (TB); or83@cornell.edu, https://orcid.org/0000-0001-8633-2302 (OR); jesseyao@cuhk.edu.hk, https://orcid.org/0000-0001-9831-5083 (Y(J)Y)

Abstract. Advances in behavioral targeting allow media firms to better monetize based on advertising. However, behavioral ad targeting requires consumer tracking, which has heightened privacy concerns among consumers and regulators. We examine how stricter privacy regulations that ban consumer tracking affect media firms' content strategies and ideological positioning. We consider a model where media firms choose their ideological positioning and advertising, whereas ideologically heterogeneous consumers select their preferred content based on both their ideology and idiosyncratic taste shocks. We compare two salient informational environments: (1) behavioral ad targeting, where perfect inference about consumers is allowed, and (2) contextual ad targeting, where consumer tracking is banned due to privacy regulations, and media firms can only infer consumer types based on their media choice. We show that privacy regulations that ban behavioral ad targeting incentivize media firms to shift toward more extreme and polarizing positioning in order to draw better inferences about consumer types, even though the shift to more extreme ideological positions can hurt both demand and consumer welfare from content consumption. Compared with the monopoly case, competition increases firms' inference motives and leads to more polarized content over a wider range of parameters due to an inferential complementarity effect arising from consumer self-selection. Our research uncovers a previously unexplored relationship between privacy and polarization, shedding light on the potential unintended consequences of privacy regulations in media markets.

1. Introduction

Digital publishers increasingly use advertising as a monetization strategy. At the core of ad-based monetization is behavioral ad targeting that creates a sustainable revenue stream for publishers and keeps online content mostly free. However, behavioral ad targeting requires the collection and use of consumer-level data, which leads to privacy concerns among consumers. According to a recent survey by Pew Research, more than 80% of U.S. adults are concerned about how companies use the data they collect from them, making the need for privacy regulation an issue with bipartisan support (McClain et al. 2023). In response to consumers' privacy concerns in the United States and globally, regulatory bodies and even some private firms have started taking actions to restrict consumer tracking online and protect consumer privacy. A few prominent examples include the California Consumer Privacy Act (CCPA) and the European Union's General Data Protection Regulation (GDPR) in the public sector and Apple's App Tracking

Transparency (ATT) and Google's Privacy Sandbox in the private sector.

Although the main intent of privacy regulations is to safeguard consumer privacy, a consistent finding from past empirical research is that privacy regulation hurts digital publishers (Goldfarb and Tucker 2011, Alcobendas et al. 2021, Johnson et al. 2023). Specifically, prior research suggests that the revenue loss due to privacy regulations is more pronounced for general interest (versus specialized) publishers, who have greater uncertainty about their consumer types without consumer tracking (Goldfarb and Tucker 2011). For example, in the absence of consumer tracking, a mainstream news website such as the Associated Press has a harder time inferring consumer types and interests in order to show them relevant ads, compared with a niche, ideologically extreme website such as Fox News, which has more precise information about their consumers. Thus, the negative impact of privacy regulation on news publishers and media firms largely

depends on their ideological positioning and content strategies.

The content-dependent impact of privacy regulation on media firms gives rise to an important question: Does privacy regulation affect media firms' ideological positioning and content strategies? To counter the loss imposed by privacy regulations, a media firm can shift from a mainstream positioning and content strategy (e.g., general news coverage) to a more specific and niche positioning and content strategy (e.g., polarizing opinion programs), thereby drawing more accurate inferences about consumers who choose their content and showing them more relevant ads. Such a shift in news content has important implications in the media landscape, especially given the rise in political polarization and media bias. Nevertheless, the extant research on privacy has focused on measuring the impact of regulations on media firms' market *outcomes* and largely ignored their equilibrium response in terms of positioning and content strategies. Focusing on media firms' equilibrium positioning strategies is particularly important given the empirical evidence that show media firms change their ideological positioning over time (Kim et al. 2022).[1]

In this paper, we bridge this gap and endogenize media firms' positioning in markets with and without privacy regulations that ban consumer tracking. We explicitly model the choice of mainstream and niche positioning using a simple Hotelling model of firm positioning and consumer demand. Specifically, we seek to answer the following questions:

1. How does privacy regulation affect the positioning and content strategies employed by media firms who monetize based on advertising in monopoly and duopoly settings?

2. What are the implications in terms of polarization? When do media firms have a greater incentive to shift to a more polarizing positioning?

3. What are the implications in terms of demand, consumers' consumption utility, and media firms' profits?

To answer these questions, we develop a model in which profit-maximizing media firms choose their ideological positioning and advertising, whereas ideologically heterogeneous readers consume the media content based on both its match with their ideological preferences and their own idiosyncratic taste shocks. In our model, media firms monetize by selling ad impressions to advertisers. As such, their profits are the product of two separate components: (1) an extensive margin that captures the *quantity of impressions* and (2) an intensive margin that captures the *quality of impressions*.

To study the impact of privacy regulations that ban consumer tracking (e.g., GDPR), we compare two different information environments: *behavioral ad targeting* and *contextual ad targeting*. The difference between the two comes from the possibility of tracking. Under *behavioral ad targeting*, media firms possess perfect knowledge of each consumer's type because they are allowed to track consumers. The *contextual ad targeting* scenario mimics the situation in the presence of privacy regulations that ban consumer tracking, where media firms are only allowed to use contextual information to target their ads to consumers. The comparison between contextual and behavioral ad targeting allows us to examine the downstream impact of privacy regulations on media firms' equilibrium choices and overall market outcomes.

Under behavioral ad targeting, because media firms can perfectly match ads to consumers on a one-to-one basis, they can achieve maximum quality of impressions regardless of consumers' media choice. As a result, the media firm's profit-maximization problem is greatly simplified. Because perfect ads-to-readers matching is possible, the profit-maximizing positioning is the one that maximizes the quantity of impressions.

Conversely, under contextual ad targeting, media firms no longer have access to perfect information about consumer types, so they need to rely solely on a single piece of information to infer consumer types: consumer's self-selection into media. In this case, media firms have the incentive to deviate from the demand-maximizing positioning (the equilibrium under behavioral ad targeting) and shift to positioning and content strategies that increase the quality of impressions. In other words, media firms may benefit from creating content with lower demand but sharper signals about consumers, thereby achieving higher profits by balancing the quantity and quality of impressions.

We first consider a monopolist media firm and examine its equilibrium strategies under both behavioral and contextual ad targeting regimes. We show that deviating from the demand-maximizing positioning and content strategy can be an equilibrium under contextual ad targeting when consumers are sufficiently ideologically differentiated and when their sensitivity to imperfect ad targeting (mismatches between ad and consumer type) is sufficiently high. In particular, we find that the deviation is more toward extreme and polarizing ideological positioning because media firms can sharpen the signal about consumer types by moving toward the extreme ends of consumers' ideological preferences. Notably, this incentive persists even when such a shift toward extreme content reduces both consumer welfare from content consumption and total demand.

We then study the duopoly case to examine the equilibrium outcomes under the two targeting regimes in a competitive environment. Studying the duopoly case is important because competition naturally incentivizes media firms to differentiate and move toward more polarizing positions. As such, it is not clear whether the inferential motives bolstered by privacy regulation are strong enough to shift the equilibrium to further

polarization under contextual ad targeting compared with behavioral ad targeting, as in the monopoly case. Our equilibrium analysis in the duopoly case reveals important insights.

First, we find that the media firms' ideological positioning in the equilibrium under contextual ad targeting is at least as polarizing as that under behavioral ad targeting, which establishes the positive link between privacy regulation and polarization in a competitive environment. Second, compared with the monopoly setting, we document that privacy regulation leads to more polarization even when consumers are less ideologically differentiated and when their sensitivity to imperfect ad targeting is lower. A key underlying mechanism we identify is an *inferential complementarity* effect between the two firms, which indicates that the presence of one firm aids the other firm's inference, making the inferential motives stronger for the media firms. Therefore, inferential complementarity, combined with regular product differentiation incentives under competition, strengthens the shift to more polarization due to privacy regulation under duopoly.

In summary, our paper makes several contributions to the literature. We study the link between privacy and polarization and highlight the possibility of an unintended consequence of privacy regulations in increasing content polarization. Specifically, we show that privacy regulation can lead to greater content polarization, even when it reduces total demand and consumer welfare from content consumption. Our finding is important as it goes against extensive media speculations and policy memos that cite hyper-personalization of online content and the lack of privacy regulation as a key contributor to the increased polarization over the past few decades (Pariser 2011). The well-documented mechanism is that lack of privacy leads to the possibility of more granular personalization, which creates echo chambers where consumers are only exposed to confirming views. Therefore, less privacy induces more polarization. In our model, we provide new insights into this problem: it is the *inability to personalize* that leads firms to move to more extreme and polarizing ideological positioning to sharpen their inferences about consumers. Thus, privacy regulation can lead to higher content polarization than cases with perfect content personalization. A key innovation of our framework is in explicitly modeling privacy as an inference problem. In particular, our inference-based framework suggests that by limiting what firms *know* about consumers, privacy regulations shift firms' focus to what they *want to know* about consumers, which creates an incentive to use content polarization as a device to draw more accurate inference. Our framework extends the literature on media bias and complements the existing mechanism for the relationship between privacy and polarization by introducing inference motives as an important determinant of

actions chosen by strategic players, which is largely ignored in the prior literature on media markets: notably, we show that privacy leads firms to polarize *beyond* a mere reflection of the underlying ideological partisanship among consumers; thus, the average consumer is exposed to *more ideologically extreme* content than they would be under perfect content personalization. Finally, our framework offers broader implications for the theoretical and empirical literature on privacy. From a game-theoretical standpoint, this modeling framework allows us to examine players' equilibrium responses to the privacy shifts. For empirical studies, this provides a framework to quantify the magnitude of information gain and loss in various data environments.

2. Literature Review

Our work relates to the literature on media markets. Early theoretical work in this domain considers the competition between broadcasters in the presence of the advertising market and shares insights into equilibrium outcomes in this market in terms of content provision and advertising strategies (Dukes and Gal-Or 2003, Gal-Or and Dukes 2003, Anderson and Coate 2005, Godes et al. 2009). A separate stream of work in this literature has examined content strategies as they relate to media bias and polarization (Mullainathan and Shleifer 2005, Gentzkow and Shapiro 2006, Xiang and Sarvary 2007). With the growth of digital news consumption, a series of recent studies have focused on the specific aspects of the digital context and examined pricing and content strategies by media markets (Johar et al. 2012; Sun and Zhu 2013; Ambrus et al. 2016; Athey et al. 2018; Berman et al. 2019; Lin 2020; Amaldoss et al. 2021, 2024; Jain and Qian 2021; Amaldoss and Du 2023; Ke et al. 2023). Our paper adds to this stream of work by studying a key aspect of the digital context: behavioral ad targeting and the possibility of privacy regulations. In particular, we study the impact of privacy regulations on media markets and examine the equilibrium outcomes in terms of ideological positioning and content strategies and their implications for media bias and polarization. The transition from behavioral targeting to contextual targeting due to privacy regulations is largely the reverse of the process of the initial adoption of behavioral targeting in the late 1990s (Novatiq 2024). The key difference between the two transitions is that the advertising choice was much less personalized following the introduction of behavioral tracking than it was before the launch of privacy regulations because real-time bidding was introduced in 2009 (Sweeney 2023). Therefore, the tradeoff between quantity and quality of impressions is a unique feature of the second transition that this paper studies.

Our work relates to the literature on behavioral targeting, personalization, and, more broadly, online advertising. With the advancements in targeting technologies in

advertising markets, a series of papers have studied the impact of targeting accuracy or advertising strategy on equilibrium market outcomes (Chen et al. 2001, Iyer et al. 2005, Levin and Milgrom 2010, Tucker and Zhang 2010, Bergemann and Bonatti 2011, Mayzlin and Shin 2011, Zhang and Katona 2012, Kuksov et al. 2013, Kamada and Kojima 2014, Amaldoss et al. 2016, Sayedi 2018, Choi and Sayedi 2019, Berman and Katona 2020, Rafieian and Yoganarasimhan 2021, Shin and Yu 2021, Lauga et al. 2022, Berman et al. 2023, Chaimanowong et al. 2023, Choi and Sayedi 2023, Fainmesser et al. 2023a, Shin and Shin 2023, Choi and Sayedi 2024, Johnson et al. 2024, Yao 2024, Ning et al. 2025, Yang et al. 2025). In the news media context, many have speculated that greater content personalization results in more polarization, citing the positive correlation between the rise of political polarization in the United States and the surge in personalized content delivery through online platforms (Pariser 2011).[2] Despite the widespread lay belief that causally connects personalization and polarization, empirical findings in this domain do not present a consistent viewpoint. For instance, the demographic groups in the United States that are least likely to use the Internet experienced the greatest increase in polarization (Boxell et al. 2017). Moreover, studies analyzing users' browsing histories reveal that, despite social media and the Internet being associated with greater ideological divergence among users, they also increase exposure to opposing views (Flaxman et al. 2016). Investigations focused on specific platforms and their personalized features offer conflicting results concerning the link between personalization and polarization. Notably, studies investigating Facebook's news feed algorithm, Google's search personalization, and YouTube have found limited evidence suggesting that personalization contributes to content bias (Bakshy et al. 2015, Ribeiro et al. 2020, Hosseinmardi et al. 2021). In our paper, we build on the theoretical literature on ad targeting and develop a model to study the impact of ad personalization on the supply of polarizing content. Our work extends this literature by providing an inference-based theoretical account that presents a more nuanced view of the link between personalization and polarization. In particular, we find that the content produced under no personalization can become *more extreme and polarized* than consumers' preexisting ideological preferences. Therefore, we highlight a case where the content under full personalization is less polarized than that under no personalization.

Our paper also relates to the literature on privacy. A vast body of theoretical work has examined different issues related to consumer identification, privacy, and information markets (Villas-Boas 1999, 2004; Taylor 2004; Acquisti and Varian 2005; Bergemann and Bonatti 2015; Bergemann et al. 2018; Choi et al. 2020; Choi and Jerath 2022; Hu et al. 2022; Iyer and Singh 2022; Momot

and Salikhov 2022; Yang 2022; Bonatti et al. 2023; Choi et al. 2023; Fainmesser et al. 2023b; Ke and Sudhir 2023; Lei et al. 2024; Miklós-Thal et al. 2024; Rhodes and Zhou 2024; Strack and Yang 2024; Yao 2025). Empirical papers in this domain have studied the impact of privacy regulations on market outcomes in different settings (Goldfarb and Tucker 2011; Johnson 2022; Johnson et al. 2023). In particular, Goldfarb and Tucker (2011) study a change in tracking and targeting regulations and document lower response rates to ads and, therefore, lower ad revenues for publishers. Notably, they demonstrate a heterogeneous effect of privacy regulation on digital publishers, with general interest publishers such as the New York Times experiencing higher revenue loss than specialized publishers such as Car and Driver Magazine. We extend this literature by endogenizing media firms' (publishers) ideological positioning and content strategies, to allow them to respond optimally to the change in privacy policies. We present a generic theoretical framework that characterizes privacy as an inference problem, in which stricter privacy policies have a negative impact on the accuracy of inference about consumers. Importantly, we identify the possibility of increased polarization as an unintended consequence of privacy regulations in digital markets.

3. Model

We theoretically characterize a market where media firms (e.g., news publishers) create content for consumers and monetize by placing ads. To reflect consumers' ideological preferences, we assume three discrete consumer types $\{0, 1/2, 1\}$, where 0 and 1 refer to the opposing ideological ends (e.g., left versus right) and $1/2$ refers to the centrist position. Consistent with consumer preferences, $\{0, 1/2, 1\}$ defines the action set for the media firm's ideological positioning or slant, which allows us to characterize the match between the consumer type and the media firm's content. The media firm's ideological positioning informs its general content strategies. As such, we refer to this decision with terms such as ideological positioning, content strategy, and slant interchangeably throughout the paper. Whereas politics is a natural application for our model, our framework encompasses broader settings with ideological preferences about issues (e.g., animal rights versus hunting websites). We unify our concept by referring to the positioning in the middle ($= 1/2$) as *mainstream* and the positioning at the two extremes ($\in \{0, 1\}$) as *niche*.[3]

We denote the consumer type by θ. We assume that the consumer type has a symmetric distribution that depends on a single parameter λ, with $1 - 2\lambda$ proportion of users having a centrist or mainstream position ($\theta = 1/2$), and λ proportion of consumers being at each extreme or niche end of consumer types ($\theta \in \{0, 1\}$). As such, we can view λ as a measure of consumer

polarization and ideological partisanship. Because we want to study media firms' incentive to polarize in equilibrium with and without privacy regulations, we focus on cases where mainstream consumers are in the minority: $\lambda \geq 1/3$. It is clear that cases with low levels of consumer polarization will lead to minimal content polarization regardless of the targeting regime. In contrast, our focal cases where $\lambda \in [1/3, 1/2]$ represent settings where content polarization can be viewed as a major possibility.[4]

We denote the media firm's ideological positioning choice by x, which comes from $\{0, 1/2, 1\}$. If the consumer chooses the media firm, an impression will be generated, and the media firm can place an ad $a \in [0,1]$ in that impression catering to different consumer types. As such, although the ideological positioning x is the same for all consumers, ad a can be targeted at the impression level depending on the information available about the consumer.[5]

We consider two types of information environment, *behavioral ad targeting* and *contextual ad targeting*. The information environment is characterized by the observability of each consumer's type, which depends on whether behavioral tracking and targeting are allowed. When *behavioral ad targeting* is allowed, we assume that the media firm knows the consumer type for each consumer on its platform, as it has rich prior information about the consumer through tracking. On the other hand, when behavioral ad targeting is banned due to privacy regulations, the media firm needs to engage in *contextual ad targeting* and infer consumer type based on the context, for example, the content consumed by the consumer.

Figure 1 presents the timeline of our game in four stages. We present the details of each stage as follows:

1. In the first stage, consumer type θ is realized. Without privacy regulation (under behavioral ad targeting), it is known by the media firm(s). With privacy regulation (under contextual ad targeting), it is the consumer's private information.

2. In the second stage, each media firm chooses an ideological positioning strategy x to maximize its expected profits. The medi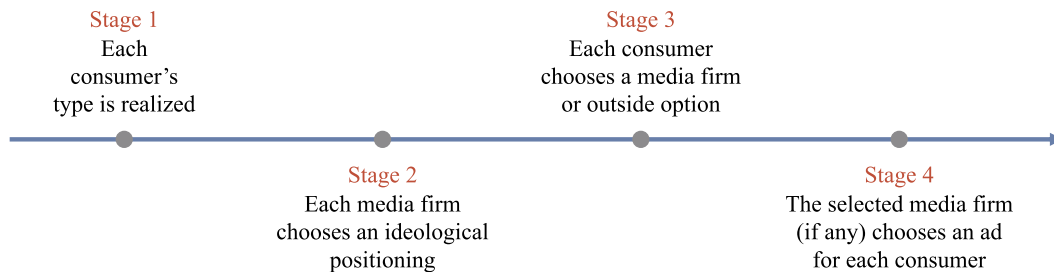a firm's choice x informs its content strategy. The media firm can choose among one mainstream position at $x = 1/2$ and two niche positions at $x = 0$ and $x = 1$. In the political context, the niche positions can be interpreted as left- and right-leaning ideological positions. The media firm's positioning choice is broadcasted to all consumers through their content. Consistent with the literature on media slant, the ideological positioning or slanting strategy is a single choice, so the firm cannot be right- and left-leaning at the same time (Mullainathan and Shleifer 2005, Gentzkow and Shapiro 2010).[6]

3. In the third stage, each consumer consumes at most one content unit.[7] We define a type θ consumer's utility from consuming the content by the media firm with ideological positioning x as follows:

$$U(x; \theta) = 1/2 - |x - \theta| + \epsilon, \qquad (1)$$

where the first term is the base utility from media consumption that is set to $1/2$, the second term is the consumer's distaste for ideological mismatch as measured by the distance of the media firm's ideological positioning from the consumer type, and the third term is an idiosyncratic error term that is firm-specific with mean zero that comes from the distribution $f(\cdot)$ such that $f(\epsilon) > 0$, $\forall \epsilon \in [-1/2, 1/2]$. Conceptually, this idiosyncratic error term captures components of a consumer's utility from consuming content by a media firm that is unrelated to their ideology θ, such as a desire to be informed or a preference for content variety.[8] Intuitively, a higher $Var(\epsilon)$ reflects a greater openness by the consumer to ideologically mismatched content. We can generalize the base utility to any v (see Section 6.1), but fixing it at $v = 1/2$ conveys the main insights and keeps the model simple. We normalize the utility from the consumer's outside option to zero. For tie-breaking rules, we assume that the consumer randomly visits one firm with equal probability if multiple firms give the consumer the same utility, and that the consumer visits the firm if indifferent between consuming the media content and choosing the outside option. The utility framework in Equation (1) implies that a consumer located at one extreme ($\theta = 0$ or 1) will never consume the content by a media firm located at the other extreme ($x = 1$ or 0), but

Figure 1. (Color online) Timeline of the Game

Stage 1
Each consumer's type is realized

Stage 2
Each media firm chooses an ideological positioning

Stage 3
Each consumer chooses a media firm or outside option

Stage 4
The selected media firm (if any) chooses an ad for each consumer

may consume the content by a mainstream media firm ($x = 1/2$) if the consumer's idiosyncratic term is sufficiently high.[9] Mainstream consumers ($\theta = 1/2$) may consume the mainstream or either niche media firms, depending on the relative values of the idiosyncratic terms.

1. In the final stage, the media firm chooses ad a if the consumer generates an impression by choosing the media firm's content. The expected value generated from showing ad a to a consumer of type θ is given by the match value function $\mathcal{M}$ as follows:

$$\mathcal{M}(a; \theta) = 1 - \gamma(a - \theta)^2, \qquad (2)$$

where the match value convexly decreases in the mismatch between ad and consumer $|a - \theta|$, and the parameter γ captures the *targeting value* and is an instrumental parameter for our analysis.[10] It is worth emphasizing that although the media firm chooses the ad, having this choice is not a requirement in our framework and this choice can be delegated to ad intermediaries (e.g., Google) who run auctions among advertisers to allocate ad impressions. We demonstrate the equivalence of our modeling framework with one that allocates ads in a competitive auction marketplace in the Online Appendix. We assume that the match value function $\mathcal{M}(\cdot; \cdot)$ is a continuous and smooth function in both a and θ. As we can see from Equation (2), the maximum value of the match value function is one. We do not restrict ourselves to the case of $\gamma \leq 1$; therefore, a large ad-consumer mismatch ($(a - \theta)^2 > 1/\gamma$) can result in negative ad profits. These could be interpreted, for instance, as the ad being repulsive to consumers (e.g., a hunting rifle ad for a vegetarian consumer), alienating them in the future.

We now define the media firm's profit maximization problem. Let $D(x)$ denote the total consumer demand for a media firm at position x. This is the *quantity* of impressions the media firm is able to generate. For each impression, the firm can place an ad and obtain ad revenue that depends on the *quality* of the impression characterized by the match value function. We assume that there is a perfectly competitive ad marketplace, so the revenue-per-impression is the same as the expected match value of the ad for a given impression.[11] Let $\mathcal{I}$ denote the information available for each impression. We define the firm's optimal profits π given information $\mathcal{I}$ as follows:

$$\pi(x, a; \mathcal{I}) = \max_{x, a} D(x) \cdot \mathbf{E}[\mathcal{M}(a; \theta) | \mathcal{I}], \qquad (3)$$

where the firm jointly maximizes content strategy x and ad a given the information $\mathcal{I}$ available.[12] The firm's contextual information about an impression is that the consumer has chosen the media firm's content x. In any event, the firm has contextual information about an impression, so we have $x \in \mathcal{I}$. Under behavioral ad

targeting, we assume that the firm also has information about the exact consumer type θ. Because we are interested in cases under *behavioral* and *contextual* ad targeting, we characterize the information available under each case as follows:

• The information under *behavioral* ad targeting is defined as $\mathcal{I}_b$ and includes the contextual information x and behavioral information θ, that is, $\mathcal{I}_b = \{x, \theta\}$.

• The information under *contextual* ad targeting is defined as $\mathcal{I}_c$ and includes the contextual information x, that is, $\mathcal{I}_c = \{x\}$.

We specifically model the information under behavioral and contextual ad targeting as extreme cases to obtain cleaner results and better understand the incentives under these two information environments. However, our framework is flexible, and one could easily consider more middle-ground cases where there is more contextual information or less behavioral information. For example, instead of a complete ban, regulations such as GDPR allow users to opt in for behavioral targeting. In such cases, firms have information about some consumers even under privacy regulations. Moreover, the probability of opting in for behavioral targeting may be correlated with user type. If mainstream consumers are more likely to opt in, then a mainstream firm will face a higher uncertainty about consumer type, whereas a niche firm will face a lower uncertainty about consumer type among consumers who do not opt in when a privacy regulation is launched. This will strengthen a firm's inference motives and, thereby, our result. In contrast, if niche consumers are more likely to opt in, then a mainstream firm will face a lower uncertainty about consumer type, whereas a niche firm will face a higher uncertainty about consumer type among consumers who do not opt in when a privacy regulation is launched. This will weaken a firm's inference motives and, thereby, our result. Empirically, most consumers do not opt in for behavioral tracking when the consent dialog is explicit and clear (Laziuk 2021, Goldberg et al. 2024). Therefore, the impact of considering the user's choice of privacy terms on the main result is small.

4. Consumer Privacy and Firm's Inference Problem

A natural way to think about consumer privacy in our setting is to examine how uncertain the firm is about the consumer's type given the information $\mathcal{I}$. We present our definition of consumer privacy as follows.

Definition 1 (Consumer Privacy). The consumer's privacy given the information $\mathcal{I}$ is defined as $Privacy(\mathcal{I})$ and is equal to the conditional variance of consumer type given information $\mathcal{I}$, that is, $Privacy(\mathcal{I}) = \mathbf{Var}(\theta | \mathcal{I})$.[13]

We can easily verify that the consumer's privacy is equal to zero under behavioral ad targeting, as $\mathcal{I}_b$ fully

reveals the consumer's type. On the other hand, the consumer's privacy under contextual ad targeting depends on the consumer's content/media choice. Some choices can reveal more about the consumer type, thereby reducing the posterior variance of type and consumer privacy. This is an important consideration for the media firm when choosing the ideological positioning and content strategy. We characterize the relationship between consumer privacy and content choice in the following definition.

Definition 2 (Privacy Reducing Choice). Consumer's content choice x is privacy-reducing over content choice x' if $Privacy(\{x\}) < Privacy(\{x'\})$.

This definition reflects the insight that a consumer who chooses Fox News or New York Times reveals more about their type than a consumer who chooses Associated Press. To see how the privacy measure enters the media firm's problem, we revisit the firm's profit maximization problem as presented in Equation (3). Because the ad choice is the final decision in the timeline, we start with the firm's optimal ad choice. For ad choice a, the firm needs to infer the consumer type given information $\mathcal{I}$. The following lemma characterizes the media firm's optimal ad choice.

Lemma 1. *The firm's optimal ad choice a^* is the posterior mean of consumer type given information $\mathcal{I}$, that is, $a^* = \mathbf{E}[\theta|\mathcal{I}]$.*

This lemma follows from the property of the variance and the quadratic specification of $\mathcal{M}$.[14] It indicates that the optimal ad a^* is the same as consumer type θ under behavioral ad targeting: $a^*(\theta) = \theta$, whereas it is equal to the mean of θ given the consumer's content choice x under contextual ad targeting: $a^*(x) = \mathbf{E}(\theta|x)$. Given the ad choice characterized in Lemma 1, we can establish the following result about the media firm's optimal ideological positioning as follows.

Corollary 1. *The media firm chooses the ideological positioning and content strategy x that maximizes $D(x) \cdot [1 - \gamma Privacy(\mathcal{I})]$.*

The corollary implies that the firm's profits are negatively correlated with consumer privacy. Under behavioral ad targeting, the firm fully observes consumer type (i.e., $Privacy(\mathcal{I}_b) = 0$), so the firm chooses the ideological positioning that maximizes its total demand:

$$\pi^b(x) = \max_x D(x), \qquad (4)$$

where $\pi^b(x)$ is shorthand for the profits under behavioral ad targeting, and the firm is only interested in maximizing the *quantity* of impressions. Under contextual ad targeting, the firm needs to infer consumer type given the consumer's self-selection into the media

firm's content, so the profit maximization problem simplifies to

$$\pi^c(x) = \max_x D(x)[1 - \gamma Privacy(\{x\})], \qquad (5)$$

where $\pi^c(x)$ is shorthand for the profits under contextual ad targeting, and $Privacy(x)$ is the privacy of consumer type under contextual ad targeting where the only information available about the consumer is their content choice x. As shown in Equation (5), the media firm has to trade off the *quantity* of impressions with the *quality* of impressions. As a consequence, the demand-maximizing ideological positioning is not necessarily the optimal one; the firm may trade off more popular positioning with a more accurate inference about the consumer. In our equilibrium analysis, we examine this inference incentive as it relates to polarization.

5. Equilibrium
5.1. Equilibrium Concept
Because of the setup of a multistage game with incomplete information, we consider pure strategy perfect Bayesian equilibrium. Specifically, when the firm chooses the advertising location in the last period, it uses Bayes' rule to update its belief about consumer type based on the consumer's equilibrium content choice.

5.2. Monopoly
5.2.1. Demand. When there is only one firm, consumers compare the utility of consuming the monopoly's content with the utility of choosing the outside option. Thus, they will consume the content if and only if the utility from content consumption is positive. The following lemma calculates consumer demand given different content choices.

Lemma 2. *The total demand for a monopolist media firm depends on their ideological positioning strategy as follows:*

a. The total demand for a monopoly's mainstream positioning, $x = 1/2$, is $1 - 2F(0)\lambda$. Among them, $1 - 2\lambda$ are mainstream consumers, $[1 - F(0)]\lambda$ are type 0 consumers, and $[1 - F(0)]\lambda$ are type 1 consumers.

b. The total demand for a monopoly's niche positioning, $x = 0$ or 1, is $1 - F(0) + [2F(0) - 1]\lambda$. Among them, λ are type $\theta = x$ consumers and $[1 - F(0)](1 - 2\lambda)$ are mainstream consumers.

This lemma highlights an important, general, and natural property. On average, a random consumer is farther from a niche position than a mainstream position. Hence, if the media firm chooses the mainstream ideological positioning, $x = 1/2$, both types of niche consumers, $\theta \in \{0, 1\}$, may consume it because they are not located too far away from the content. In contrast, if the firm chooses the niche ideological positioning, say $x = 0$, mainstream consumers may consume it, whereas type $\theta = 1$ consumers will not, independent of their

idiosyncratic taste shocks, because they are too ideologically distant from the content. This endogenous selection leads to an asymmetry in the firm's inference ability. By choosing niche ideological positioning $x = 0$, the firm knows for sure that it does not attract any type $\theta = 1$ consumers and can better infer the consumer's type by Bayesian updating.

5.2.2. Media Firm's Ideological Positioning. The monopolist's positioning strategy is straightforward under behavioral ad targeting. It chooses the ideological positioning that maximizes total demand because there is no inference problem, and profit simply equals total demand, as indicated in Equation (4). The following proposition characterizes the monopolist equilibrium under behavioral ad targeting.

Proposition 1. *Under behavioral ad targeting, the monopolist media firm chooses the mainstream ideological positioning if* $[4F(0) - 1]\lambda < F(0)$ *and niche ideological positioning if* $[4F(0) - 1]\lambda > F(0)$.

This proposition characterizes the equilibrium strategy under the general case of behavioral ad targeting. The optimality condition for the mainstream positioning can easily be satisfied. For example, if we make the common assumption that the distribution of the idiosyncratic term is symmetric, we find that the mainstream positioning is always the equilibrium strategy. This is because $F(0) = 1/2$ under a symmetric distribution, which implies that for all possible values of λ, mainstream positioning is the equilibrium strategy.

As discussed earlier, the firm's perfect inference under behavioral ad targeting implies that the firm does not strategically consider inference. In contrast, under contextual ad targeting, it relies on Bayesian updating to infer consumer type and considers inference strategically. As such, it may have an incentive to deviate from the demand-maximizing strategy (i.e., equilibrium under behavioral ad targeting) and choose an ideological positioning that allows more accurate ad targeting. Hence, the question is under which strategy the firm knows more about consumer type in an impression. The following lemma helps characterize the decision-making tradeoff for the media firm.

Lemma 3. *The consumer's choice of niche content is privacy-reducing over the choice of mainstream content, that is, Privacy($\{1/2\}$) > Privacy($\{x\}$), such that $x \in \{0, 1\}$.*

This lemma indicates that a mainstream positioning leads to a lower intensive margin (ad revenue) on each impression, because the firm will have greater uncertainty about the consumer type given their content choice. Therefore, there is a natural incentive for media firms to move to more niche ideological positions. If the niche positioning is demand-maximizing and thus the equilibrium under behavioral ad targeting (when

$[4F(0) - 1]\lambda > F(0)$), Lemma 3 implies that it will be the equilibrium under contextual ad targeting, because niche strategy does better on both elements in the expected profit equation in Equation (5): *quantity* and *quality* of impressions.

When the demand-maximizing strategy is mainstream ideological positioning, the tradeoff between quantity and quality of impressions becomes more interesting. Naturally, we expect the equilibrium to depend on the value of targeting that is captured by the parameter γ. The following proposition characterizes this relationship.

Proposition 2. *Under contextual ad targeting, the monopolist media firm's equilibrium strategy depends on the values of λ and γ as follows:*

a. *If* $[4F(0) - 1]\lambda > F(0)$, *the media firm chooses niche ideological positioning in equilibrium.*

b. *If* $[4F(0) - 1]\lambda < F(0)$, *there exists $\gamma^m > 0$ for any value of λ such that the media firm chooses niche ideological positioning if $\gamma > \gamma^m$ and mainstream ideological positioning if $\gamma < \gamma^m$. The threshold γ^m is decreasing in λ.*

Figure 2 combines the equilibrium under behavioral and contextual ad targeting and presents three regions. We find that equilibrium strategies are only different in region B (shaded) where $[4F(0) - 1]\lambda < F(0)$ and $\gamma > \gamma^m$. In this case, the monopolist media firm chooses a niche ideological positioning in equilibrium under contextual ad targeting, but a mainstream positioning under behavioral ad targeting. Although the media firm generates fewer impressions by choosing a niche ideological

Figure 2. (Color online) Equilibrium Regions in Monopoly Under Behavioral and Contextual Ad Targeting

Notes. This figure shows equilibrium regions when $F(0) = 9/16$. Region A shows the cases where mainstream positioning is the equilibrium under both targeting regimes. Region B (shaded) illustrates cases where privacy regulation leads to higher polarization because the firm chooses niche positioning under contextual ad targeting, but a mainstream positioning under behavioral ad targeting. Region C shows cases where the firm chooses niche positioning under both targeting regimes.

positioning, it faces lower uncertainty about consumer type in these impressions. This lower uncertainty translates into a higher revenue-per-impression, which makes up for the lower demand when the value of targeting is sufficiently high, that is, $\gamma > \gamma^m$. Together, the ideological positioning that the media firm chooses under contextual ad targeting is at least as polarizing as that under behavioral ad targeting.

From a privacy standpoint, our analysis shows that a privacy regulation that bans consumer tracking can lead to increased polarization and partisanship in the supply side. Intuitively, this is because consumer tracking better separates the media firm's ideological positioning choice from its advertising choice. Moreover, our results shed light on the well-studied link between consumers' ideological partisanship (demand) and content polarization (supply), where the prior literature finds a natural incentive for media firms to slant their content (Gentzkow and Shapiro 2010). We revisit this link with and without privacy regulations and find that when consumers are more ideologically partisan (higher λ), privacy regulation can change the media firm's equilibrium strategy from moderate (mainstream) to partisan (niche) for a wider range of γ, because the threshold γ^m is decreasing in λ. This is a natural comparative static: When more consumers have strong ideological positions, the media firm is more likely to pander to them, even for lower values of targeting. Thus, we find that privacy regulation strengthens the link between consumers' ideological partisanship (demand) and content polarization (supply).

5.2.3. Consumer Utility from Content Consumption.
We have shown that the monopolist media firm may trade off the total demand for a more accurate inference when behavioral tracking is banned. In this section, we ask the following question: how does the switch from mainstream to niche content strategy affect consumers' utility from content consumption? Intuitively, niche consumers receive a higher utility from consuming niche content, whereas mainstream consumers receive a higher utility from consuming mainstream content. Hence, niche content strategy leads to lower consumer utility from content consumption if most consumers are mainstream (low λ), whereas it leads to higher consumer utility from content consumption if most consumers are niche (high λ). We formalize this insight in the following proposition.

Proposition 3. *Niche content strategy leads to lower consumer utility from content consumption if and only if $\lambda < \hat{\lambda}$:= $\{1 - 2[1 - F(0)]\mathbf{E}[\epsilon|\epsilon > 0]\}/\{3 - 8[1 - F(0)]\mathbf{E}[\epsilon|\epsilon > 0]\}$.*

We can now combine this proposition with Propositions 1 and 2 and examine when the equilibrium strategy under contextual ad targeting leads to higher or lower consumer utility from content consumption compared with the equilibrium under behavioral ad targeting. This is an important comparison from a privacy policy standpoint because it helps us understand a key welfare consequence of privacy regulations. To illustrate this aspect, we need to focus on where the equilibrium strategies are different with and without privacy regulation ($[4F(0) - 1]\lambda < F(0)$ and $\gamma > \gamma^m$) and examine under what condition privacy regulation has unintended consequences in terms of consumer welfare from content consumption. The following corollary highlights this insight.

Corollary 2. *Consumer utility from content consumption is lower under contextual ad targeting than under behavioral ad targeting if $\lambda < \hat{\lambda}$, $[4F(0) - 1]\lambda < F(0)$, and $\gamma > \gamma^m$.*

The cutoff threshold $\hat{\lambda}$ determines the region where privacy regulation has an unintended consequence in terms of consumer welfare. This cutoff depends critically on the distribution of the idiosyncratic term. In particular, $\hat{\lambda}$ is increasing in $\mathbf{E}(\epsilon|\epsilon > 0)$.[15] Therefore, understanding $\mathbf{E}(\epsilon|\epsilon > 0)$ is essential to the interpretation of our results.

In general, a distribution of ϵ that is concentrated around 0 results in low values of $\mathbf{E}(\epsilon|\epsilon > 0)$, whereas the opposite is true for a bimodal distribution that concentrates mass near the extremes $-1/2$ and $1/2$, for which $\mathbf{E}(\epsilon|\epsilon > 0) \to 1/2$. As such, $\mathbf{E}(\epsilon|\epsilon > 0)$ has a close relationship with the variance of the idiosyncratic term ϵ, which captures consumers' media preferences that are unexplained by ideology. In that sense, the higher the variance of the idiosyncratic term, the more open the consumer is to consuming content farther from their own ideology. This is an important construct because many behavioral nudges to encourage people to read from other sides attempt to reduce the relative importance of ideology in consumer utility.

Interestingly, our results suggest that as consumers' openness to ideologically different content increases (higher variance of ϵ), the region where privacy regulation reduces consumer welfare from content consumption grows (higher $\hat{\lambda}$). Focusing on the two extreme cases is illuminating. First, if there is no variance in the idiosyncratic term ($\epsilon = 0$), the cutoff $\hat{\lambda}$ will be equal to $1/3$, implying that the privacy regulation always increases consumer welfare whenever it changes the equilibrium under behavioral ad targeting. In contrast, in the limit case of a bimodal ϵ where $\mathbf{E}(\epsilon|\epsilon > 0) \to 1/2$, we have $\hat{\lambda} = 1/2$, which indicates that the equilibrium shift caused by privacy regulation reduces consumer welfare from content consumption. We illustrate this point in Figure 3, where we use three different error distributions by varying the shape parameters of the Beta distribution, upon normalizing its expected value to zero. As shown in each of the three panels in Figure 3, as the relative importance of ideology in consumer preferences decreases through changing the variance of the

Figure 3. (Color online) Impact of Privacy Regulation on Consumer Utility from Content Consumption

(a) Bimodal ϵ:
$\epsilon \sim Beta(0.1, 0.1) - 1/2$

(b) Uniform ϵ:
$\epsilon \sim Beta(1, 1) - 1/2$

(c) Bell-shaped ϵ:
$\epsilon \sim Beta(10, 10) - 1/2$

Notes. The shaded region with the plus sign shows the region where consumer utility from content consumption is higher under contextual ad targeting than behavioral ad targeting. The striped region with the minus sign shows the region where the equilibrium under contextual ad targeting leads to lower consumer utility from content consumption. The blank region shows the area where the consumer utility is the same under both behavioral and contextual ad targeting.

idiosyncratic term, the region that reduces consumer utility from content consumption expands.

In sum, our analysis shows several important welfare implications. First, we show that there always exists a parameter range such that privacy regulation leads to a lower consumer utility from content consumption.[16] This implies that banning behavioral tracking may drive the monopolist media firm to switch from mainstream to niche ideological positioning, even when doing so reduces both total demand and consumer welfare from content consumption. Second, we show that the negative impact on consumer welfare from content consumption is more pronounced when ideology plays a smaller role in determining consumers' media preferences, a case that is often the target of depolarization efforts.

Before concluding our study of the monopoly case, we emphasize that the negative consequence of privacy regulation highlighted in Proposition 3 and Corollary 2 depends solely on consumers' utility from content consumption and does not account for additional utility such as consumer's disutility from polarization, from decreased privacy, and from potentially being exposed to offensive (that is, heavily mismatched) ads. In addition, the social welfare may suffer from a higher level of polarization. Because the impact of privacy regulation on polarization and on the consumer's utility from content consumption is of sufficient importance and interest, we abstract away from the above considerations in this paper. Because of the above reasons, increasing content polarization per se may be viewed as an unintended consequence of privacy regulations. In a richer model where one takes into account the additional utility terms, privacy regulation can hurt consumer or social welfare

by increasing the level of polarization and decreasing consumer privacy even if its direct impact on consumer's utility from content consumption is nonnegative.

5.3. Duopoly

We now introduce competition in a duopoly setting. The impact of competition in our setting is both a priori unclear and conceptually interesting: Are the two firms' inference motives complements or substitutes? Further, given that competition generally leads to an increase in differentiation, is it still true that contextual ad targeting leads to *more* content polarization compared with behavioral ad targeting? Moreover, provided that an increase in differentiation occurs following a ban on individual data tracking, what are its implications for consumer demand and utility from media consumption? To answer these questions, we consider a duopoly setting in this section. Note that when there are multiple firms at the same position, the consumer's idiosyncratic term ϵ for each firm's content is drawn independently from the distribution $f(\cdot)$, which captures a taste shock for consuming content by each media firm.

5.3.1. Demand. In contrast to the monopoly case, consumers compare the utility of consuming a firm's content not only with the outside option but also with the utility of consuming the other firm's content. Suppose that the first media firm chooses x_1 ideological positioning and the second one chooses x_2 ideological positioning. Let (x_1, x_2) denote firms' positioning strategies. Without loss of generality, we assume that $x_1 \leq x_2$. By symmetry, $(0,0)$ is equivalent to $(1,1)$, and $(0,1/2)$ is equivalent to $(1/2,1)$. Hence, we only need to consider $(0,0)$, $(0,1/2)$, $(0,1)$, and $(1/2,1/2)$. One can see that $(0,1)$

dominates $(0,0)$, so $(0,0)$ will never be an equilibrium. Hence, we need to only consider three possible equilibria: $(0,1)$, $(1/2,1/2)$ and $(0,1/2)$. Lemma 4 fully characterizes consumer demand under each possible equilibrium strategies for the two firms as follows.

Lemma 4. *The media firms' demand under three possible equilibria* $(0,1)$, $(1/2,1/2)$, *and* $(0,1/2)$ *is characterized as follows:*

a. *Under* $(0,1)$ *strategy profile, Firm 1's total demand is* $\lambda + [1 - F(0)^2](1 - 2\lambda)/2$. *Among them,* λ *are type 0 consumers and* $[1 - F(0)^2](1 - 2\lambda)/2$ *are mainstream consumers. Firm 2's total demand is* $\lambda + [1 - F(0)^2](1 - 2\lambda)/2$. *Among them,* λ *are type 1 consumers and* $[1 - F(0)^2](1 - 2\lambda)/2$ *are mainstream consumers.*

b. *Under* $(1/2,1/2)$ *strategy profile, each firm's total demand is* $(1 - 2\lambda)/2 + [1 - F(0)^2]\lambda$. *Among them,* $(1 - 2\lambda)/2$ *are mainstream consumers,* $[1 - F(0)^2]\lambda/2$ *are type 0 consumers, and* $[1 - F(0)^2]\lambda/2$ *are type 1 consumers.*

c. *Under* $(0,1/2)$ *strategy profile, Firm 1's total demand is* $\lambda[1 - \int_{-1/2}^{0} F(\epsilon_0)f(\epsilon + 1/2)d\epsilon] + (1 - 2\lambda)\int_{-1/2}^{0} F(\epsilon)f(\epsilon + 1/2)d\epsilon$. *Among them,* $\lambda[1 - \int_{-1/2}^{0} F(\epsilon)f(\epsilon + 1/2)d\epsilon]$ *are type 0 consumers and* $(1 - 2\lambda)\int_{-1/2}^{0} F(\epsilon)f(\epsilon + 1/2)d\epsilon$ *are mainstream consumers. Firm 2's total demand is* $\lambda\int_{-1/2}^{0} F(\epsilon)f(\epsilon + 1/2)d\epsilon + (1 - 2\lambda)[1 - \int_{-1/2}^{0} F(\epsilon)f(\epsilon + 1/2)d\epsilon] + [1 - F(0)]\lambda$. *Among them,* $\lambda\int_{-1/2}^{0} F(\epsilon)f(\epsilon + 1/2)d\epsilon$ *are type 0 consumers,* $(1 - 2\lambda)[1 - \int_{-1/2}^{0} F(\epsilon)f(\epsilon + 1/2)d\epsilon]$ *are type 1/2 consumers, and* $[1 - F(0)]\lambda$ *are type 1 consumers.*

Even without inference motivation (i.e., in the case of behavioral ad targeting), competing firms have an incentive to choose a more polarized content strategy to increase differentiation and soften the competition, which increases the total demand. Therefore, it is, in some sense, harder for them to choose a *more* polarized content strategy (compared with the baseline level of differentiation) when the information environment shifts from behavioral ad targeting to contextual ad targeting. Nevertheless, the next section shows that the inference motivation becomes stronger with competition.

5.3.2. Media Firms' Ideological Positioning. For tractability, we make an additional assumption that $\epsilon \sim U[-1/2, 1/2]$, which simplifies the demand function characterized in Lemma 4. The following propositions summarize the equilibria under behavioral and contextual ad targeting.

Proposition 4. *Suppose* $\epsilon \sim U[-1/2, 1/2]$. *Under behavioral ad targeting, the equilibrium is* $(1/2,1/2)$ *if* $\lambda < \lambda_1$, $(1/2,1/2)$ *or* $(0,1)$ *if* $\lambda_1 < \lambda < \lambda_2$, *and* $(0,1)$ *if* $\lambda > \lambda_2$, *where* $1/3 < \lambda_1 = 4/11 < \lambda_2 = 3/7$.

Proposition 5. *Suppose* $\epsilon \sim U[-1/2, 1/2]$. *The equilibria are (weakly) more polarized under contextual ad targeting than under behavioral ad targeting.*[17] *Table 1 summarizes*

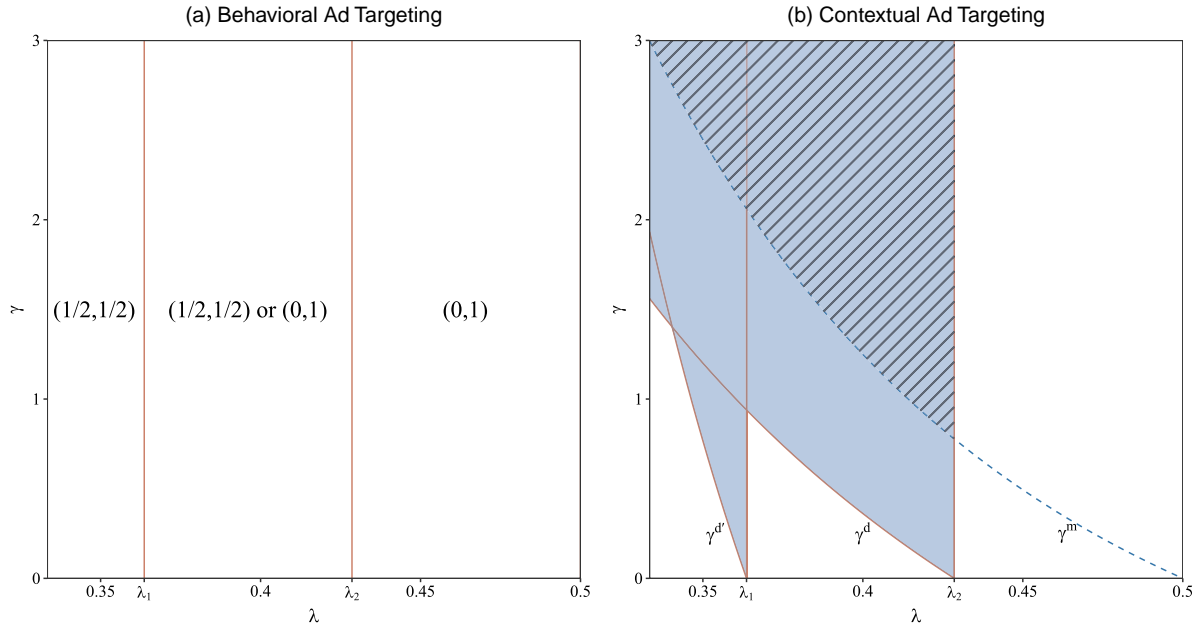**Table 1.** Equilibria Under Behavioral and Contextual Ad Targeting in Duopoly

	Equilibrium under	
Condition	Behavioral ad targeting	Contextual ad targeting
$1/3 < \lambda < \lambda_0$ & $\gamma^d < \gamma < \gamma^{d'}$	$(1/2,1/2)$	$(0,1/2)$
$\lambda_0 < \lambda < \lambda_1$ & $\gamma^{d'} < \gamma < \gamma^d$	$(1/2,1/2)$	$(1/2,1/2)$ or $(0,1)$
$1/3 < \lambda < \lambda_1$ & $\gamma > \max\{\gamma^d, \gamma^{d'}\}$	$(1/2,1/2)$	$(0,1)$
$\lambda_1 < \lambda < \lambda_2$ & $\gamma > \gamma^d$	$(1/2,1/2)$ or $(0,1)$	$(0,1)$

cases where the media firms choose more polarized content strategies under contextual ad targeting than under behavioral ad targeting, where $1/3 < \lambda_0 < \lambda_1$, $\gamma^d > 0$ *for* $\lambda < \lambda_2$, *and* $\gamma^{d'} > 0$ *for* $\lambda < \lambda_1$.[18]

Figure 5 illustrates media firms' content strategies in duopoly under behavioral (left) and contextual (right) ad targeting. The content strategies are the same under both behavioral and contextual ad targeting in the blank region, whereas the extent of media polarization is larger under contextual ad targeting than under behavioral ad targeting in the solid and striped-shaded regions.

When $\lambda > \lambda_2$, there are enough niche consumers (at least 6/7th of the total) that $(0,1)$ is the equilibrium in duopoly even under behavioral ad targeting. As such, there is no room for the duopoly to be more polarized if the advertising environment switches from behavioral ad targeting to contextual ad targeting. Now consider the interesting case in which $\lambda < \lambda_2$, so that at least 1/7th of consumers are mainstream. In both monopoly and duopoly, contextual ad targeting leads to more content polarization when $\gamma > \gamma^m$ (the striped shaded region in Figure 4(b)). In addition, contextual ad targeting leads to more polarization when $\min\{\gamma^d, \gamma^{d'}\} < \gamma < \gamma^m$ and $\lambda < \lambda_1$ or when $\gamma^d < \gamma < \gamma^m$ and $\lambda_1 < \lambda < \lambda_2$ (the solid shaded region in Figure 4(b)) under duopoly but not under monopoly. In particular, when $1/3 < \lambda < \lambda_1$ and $\gamma > \max\{\gamma^d, \gamma^{d'}\}$, equilibria across the two regimes display opposite properties: Equilibrium under behavioral ad targeting features both firms choosing the mainstream ideological positioning, whereas equilibrium under contextual ad targeting features the two firms picking opposite niche ideological positions. Thus, privacy regulation leads firms to switch from "underproviding" to "overproviding" (compared with the underlying consumer ideology distribution) highly partisan positions.

Therefore, even *conditional on differentiation motives due to competition*, duopolistic firms have a *stronger* incentive than a monopoly to polarize due to inference motives. In other words, compared with a monopolist, duopolistic firms choose a more polarized strategy under contextual rather than behavioral ad targeting in a wider range of parameters, as long as the competition is not strong enough that they already have chosen the most polarized strategy under behavioral ad targeting.

Figure 4. (Color online) Media Firms' Equilibrium Content Strategies in Duopoly When $\epsilon \sim U[-1/2, 1/2]$

Notes. The blank region shows areas where the equilibrium is the same under behavioral and contextual ad targeting in duopoly. The shaded area shows the total area where the equilibrium is more polarized under contextual ad targeting than behavioral ad targeting. The striped region shows the part of the shaded area where the equilibrium under contextual ad targeting is more polarizing than the equilibrium under behavioral ad targeting in monopoly too.

The underlying mechanism is the following. On the one hand, competing media firms directly cannibalize each other's demand if they choose the mainstream content strategy. In contrast, they can cover different niche consumers by choosing the opposite ideological positions and thereby soften the competition. As a result, the

benefit of higher total demand in the market by choosing mainstream content strategy is lower in the duopoly case than in the monopoly case. On the other hand, the benefit of a more accurate inference is not affected by competition: A monopolistic firm and a duopolistic firm extract the same amount of surplus in the advertising

Figure 5. (Color online) Impact of Privacy Regulation on Consumer Utility from Content Consumption (Duopoly Case)
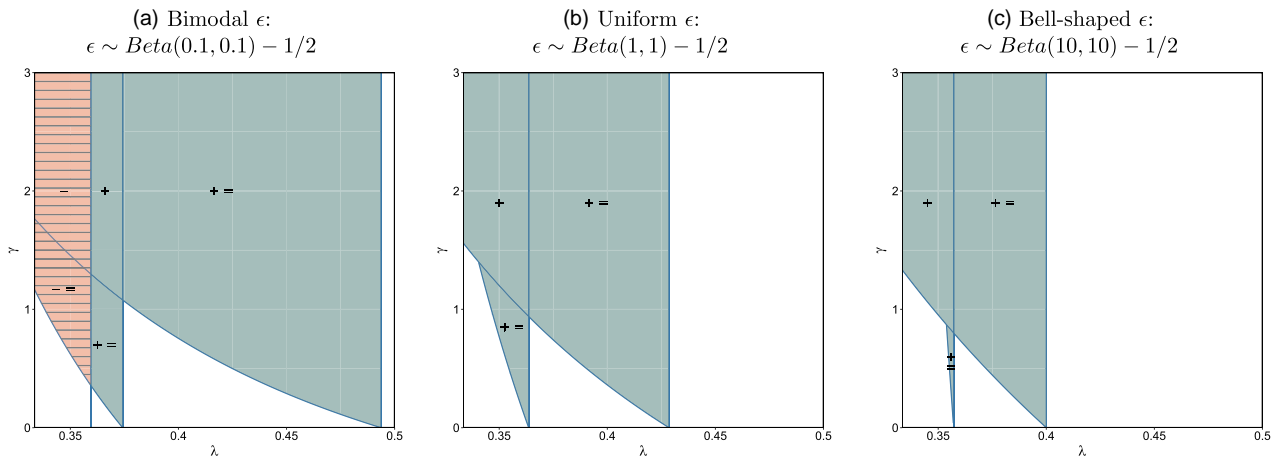
Notes. The shaded region with the plus sign shows the region where consumer utility from content consumption is always higher under contextual ad targeting than behavioral ad targeting. The shaded region with the plus-equal sign shows the region where consumer utility from content consumption is higher under contextual ad targeting than behavioral ad targeting in some equilibria and identical in other equilibria. The striped region with the minus sign shows the region where the equilibrium under contextual ad targeting leads to lower consumer utility from content consumption. The striped region with the minus-equal sign shows the region where the equilibrium under contextual ad targeting leads to lower consumer utility from content consumption under some equilibria and leads to identical consumer utility from content consumption under other equilibria. The blank region shows the area where the consumer utility is the same under both behavioral and contextual ad targeting.

market for a given level of consumer privacy. In addition, by choosing niche ideological positioning, a duopolistic firm can obtain a more accurate inference than a monopoly due to an *inferential complementarity* effect. The presence of one firm aids the other firm's inference because competition improves the selection of consumers: On average, consumers of each firm are more ideologically aligned with that firm, compared with the monopoly case. For instance, in the $(0, 1)$ case, both firms lose relatively more demand from consumers of opposite ideologies, thus decreasing the variance in consumer types for each firm and increasing profits per impression in the contextual ad targeting regime.[19] Therefore, firms lean more toward quality in the quantity-quality trade-off under a duopoly than under a monopoly.

The fact that an increase in privacy results in an increase in equilibrium content polarization is interesting in light of the well-discussed link between content personalization (that is, lack of privacy) and polarization. It is important to stress that, according to this view, personalized content does not *polarize* consumers; rather, it simply *matches* their prior ideological positions. In this case, one could say that content positioning is the *consequence*, rather than the *cause*, of polarization.

The drivers of Proposition 5 are quite different: The desire to make precise inferences to increase advertising profits pushes two competing firms away from the standard Hotelling equilibrium, $(1/2, 1/2)$ and toward $(0, 1)$ that is more ideologically polarized than the underlying distribution of the consumers' ideological preferences for any value of $\lambda < 1/2$. Therefore, more privacy leads to more content polarization on the supply side. Content polarization is driven by firms' desire to monetize ads and not just by consumer ideology; in fact, content distribution often becomes more polarized than consumers.

5.3.3. Consumer Utility from Content Consumption.
As we have shown in the monopoly case, a more polarized content strategy may lead to lower consumer utility from content consumption. However, with two firms, polarization becomes desirable for consumers. The next proposition shows that consumer utility from content consumption always increases in the extent of polarization under duopoly.

Proposition 6. *For any $\lambda \in (1/3, 1/2]$ and $\epsilon \sim U[-1/2, 1/2]$, a more polarizing equilibrium in duopoly leads to a higher consumer utility from consumption:* $(0, 1)$ *leads to higher consumer utility from content consumption than* $(0, 1/2)$, *and* $(0, 1/2)$ *leads to higher consumer utility from content consumption than* $(1/2, 1/2)$.

This proposition helps us examine the implications of a privacy regulation that enforces a move from behavioral to contextual ad targeting. As shown in Proposition 5, the equilibrium under contextual ad targeting is at least as polarizing as the equilibrium under behavioral

ad targeting for the entire parameter space. Thus, we can write the following corollary.

Corollary 3. *For any $\lambda \in (1/3, 1/2)$ and $\epsilon \sim U[-1/2, 1/2]$, privacy regulation leads to higher consumer utility from content consumption in a duopoly.*

This corollary shows that a privacy regulation in the duopoly context does not have unintended consequences in terms of the content consumption aspect of consumer welfare when ϵ is uniformly distributed. This is different from our findings in the monopoly setting, which shows that the shift to greater content polarization under privacy regulation can reduce consumers' utility from content consumption under uniform distribution. Intuitively, our findings in the duopoly setting indicate that consumers enjoy higher content polarization, which is expected given that we consider cases where consumers are polarized to begin with ($\lambda > 1/3$). We further run numerical simulations for other distributions of the idiosyncratic term, $\epsilon \sim Beta(0.1, 0.1) - 1/2$ and $\epsilon \sim Beta(10, 10) - 1/2$, and present the results in Figure 5. Compared with Figure 3 in the monopoly case, we find similar patterns that as consumers' content choice becomes less revealing of their ad content preference (higher variance of ϵ), the region where privacy regulation reduces consumer welfare from content consumption grows. Notably, we replicate the finding under monopoly that the region where privacy regulation reduces consumer utility from content consumption grows as consumers' openness to ideologically different content increases (higher variance of ϵ). In addition, conditional on the variance of the idiosyncratic term, privacy regulation hurts consumers' utility from content consumption less when there is competition. In summary, we find that compared with the monopoly case, in a competitive environment, privacy regulation helps consumers receive higher utility from content consumption, although at the expense of an increase in content polarization.

6. Extensions
6.1. General Base Utility
Throughout our exposition thus far, the consumer's base utility from media consumption was fixed at $1/2$. This allows us to convey all the main insights about the inference role of content strategy and the quantity-quality tradeoff while keeping the model simple. Nevertheless, it is interesting to study whether and how the incentives to opt for mainstream or niche positioning differ for firms of different quality. In this extension, we generalize the base utility to v and show the robustness of our findings. In addition, we demonstrate how the size of the base utility strengthens/weakens the firm's inference incentive.

Proposition 7. *Suppose $\epsilon \sim U[-1/2, 1/2]$. There exists $\underline{v} \le 7/25, \overline{v} \ge 5/6, \gamma^m > 0, \gamma^{d''}, \lambda_1 < \lambda_2$ and $\overline{\lambda}(v) \in (1/3, 1/2]$ such that the following holds if $\underline{v} < v < \overline{v}$:*

a. *(Monopoly) The firm chooses mainstream content under behavioral ad targeting and niche content under contextual ad targeting if $1/3 < \lambda < \overline{\lambda}(v)$ and $\gamma > \gamma^m$. The threshold γ^m decreases in λ and increases in v. In addition, there exists $\hat{\lambda}(v) > 1/3$ such that niche content strategy leads to lower consumer welfare if $\lambda < \hat{\lambda}(v)$.*

b. *(Duopoly) The equilibrium under behavioral ad targeting is $(1/2,1/2)$ if $\lambda < \lambda_1$, $(1/2,1/2)$ or $(0,1)$ if $\lambda_1 < \lambda < \lambda_2$, and $(0,1)$ if $\lambda > \lambda_2$. The equilibrium under contextual ad targeting is more polarized than that under behavioral ad targeting if $\gamma > \gamma^{d''}$ and $1/3 < \lambda < \lambda_2$. Moreover, $\gamma^{d''} < \gamma^m$ for all $\lambda \in (1/3, \lambda_2)$.*

As we can see, for an interval range of the base utility that includes $1/2$, we recover all the main insights from Propositions 1 to 5 in the main model. It is still the case that the media firm's ideological positioning has an inference role for the firm under contextual ad targeting, that a media firm will choose more polarized content in equilibrium to increase profit per consumer despite the lower total demand and potentially lower consumer welfare and that competition strengthens firms' inference incentive.

In the main model, a consumer located at one extreme will never consume the content located at the other extreme. In contrast, a consumer located at one extreme may consume the opposite content when the base utility v is larger than $1/2$. In particular, a type 0 consumer will obtain a positive utility from consuming $x = 1$ content with probability $\mathbf{P}(v - |x - \theta| + \epsilon > 0) = 1/3$ if $v = 5/6$. Therefore, our extension to the general base utility establishes the robustness of our results to settings where a consumer located at one extreme may consume content from the opposite extreme.

The general base utility setup provides additional insight into how the base utility affects the firm's inference incentive. The firm has a smaller incentive to improve its inference by choosing more polarized ideological positioning as the base utility increases because γ^m increases in v. The base utility can be viewed as the vertical quality of the media firm, whereas the match between consumer type and media/ad location is horizontal. When the vertical quality is high, the media firm attracts consumers with ideological preferences far from the firm's position. As a result, the horizontal match quality becomes less important and it becomes harder for the firm to distinguish different consumer types based on their self-selection. The firm is more incentivized to favor quantity in the quantity-quality tradeoff by choosing mainstream content with a larger total demand.

In contrast, when the base utility is low, consumers are more selective in their consumption decisions. Mainstream consumers will not consume niche content that has a low vertical quality and poor horizontal match unless the random shock is very high (a large idiosyncratic term). Therefore, by choosing the more polarized niche content strategy, the firm can screen consumer types very accurately and thus obtain a high profit per consumer by offering well-matched ads to most consumers. In this case, the firm has a stronger incentive to favor quality in the quantity-quality tradeoff by choosing a niche content strategy.

In summary, our extension to the general base utility shares important insights into the link between privacy and content polarization. In particular, we find that the impact of privacy regulation on content polarization is stronger when the base utility is lower. On the other hand, the link becomes weaker when the base utility is higher. Thus, policymakers can focus their attention on policies that incentivize firms to produce a higher vertical content quality.

6.2. Consumer Utility from Advertising

In the main model, consumers derive utility solely from media consumption. This assumption can be justified by two arguments. First, it is reasonable to expect that the consumer places an overwhelming majority of the weight on the media content rather than on the advertising content. Second, taking into account the utility of future ads would require the consumer to be forward-looking, which is not always the case.

In this extension, we allow consumers to gain more disutility from worse-matched ads and to take it into account in their content choices.[20] We show that our qualitative results still hold as long as consumers place the majority of the weight on the media content rather than on the advertising content, an assumption that has been empirically validated in recent studies (Brynjolfsson et al. 2024). Formally, we consider the following utility specification:

$$U(x, \theta) = 1/2 - |x - \theta| - \mathbf{E}[f(a, \theta)] + \epsilon,$$

where a is the advertising choice. Analogously to the rest of our analysis, under behavioral ad targeting the firm offers personalized ads to each consumer, $a = \theta$, whereas under contextual ad targeting the ads can only depend on the media content choice, $a = \mathbf{E}_x(\theta)$, which is the expected consumer type among consumers who consume the media content. We make the following assumptions on f to reflect the idea that the consumer prefers better-matched ads but places the majority of the weight on the media content.

Assumption 1. *The advertising disutility function satisfies the following conditions:*
(1) $f(a, \theta)$ *increases in* $|a - \theta|$; (2) $f(\theta, \theta) = 0$; (3) $f(a, \theta) \le M$.

In this assumption, M represents an upper bound for ad disutility. The empirical literature on ad avoidance suggests that the advertising volume leads to ad avoidance, but the evidence remains limited on the ad content

playing a major role in driving content consumption (Wilbur 2016, Brynjolfsson et al. 2024). Hence, with the same level of advertising volume, it is reasonable to expect that the upper bound M is small relative to $|x - \theta|$. To understand why we need this assumption, consider an extreme case in which consumers place a very high weight on the advertising disutility term. Suppose the firm chooses niche content $x = 0$. None of the type 0 consumers consume the content, whereas mainstream consumers consume it as long as their idiosyncratic term is positive. Under this unrealistic strategy profile, the optimal ad choice is 1/2 because only mainstream consumers consume the media content in equilibrium. The expected disutility term is zero for mainstream consumers but large for niche consumers. The difference in the advertising disutility term more than compensates the difference in the media consumption term, $1/2 - |x - \theta| + \epsilon$. Such strategic implications can overrule our mechanism in the main model.

Under behavioral ad targeting, the equilibrium is the same as the main model because the firm can perfectly infer consumer type and thus show perfectly matched ads. The case of contextual targeting is more complex. The reason for this is that under contextual ad targeting, each consumer's strategy now depends not only on the firm's strategy and the consumer's own type but also on the other consumers' actions. The following proposition shows that our main qualitative insight from earlier analysis holds.

Proposition 8. *Suppose $\epsilon \sim U[-1/2, 1/2]$. There exists $\hat{M} > 0$ such that the consumer's choice of niche content is privacy-reducing over the choice of mainstream content and that all the main insights of the monopoly case hold qualitatively if $M \leq \hat{M}$.*

6.3. General Oligopolistic Case (More Than Two Firms)

In this section, we extend the monopoly and duopoly model to a case with a fixed set of media firms. Consistent with the duopoly case, we consider firm-specific idiosyncratic error terms that are independent and come from $U[-1/2, 1/2]$. Suppose that we have $K > 2$ firms. In the Online Appendix, we derive the closed-form conditional probabilities of consumer's content choice given type, for any given number of firms at each position (N_0, M, N_1). We note that the signal becomes exponentially sharp from the content choice as the number of firms offering content x increases. In particular, as shown in the Online Appendix, with L firms at position x, the probability that a consumer of type $\theta = x$ chooses a media firm at another position is upper bounded by $(1/2)^L$. This indicates that with a very large number of firms at each position, the amount of inference firms can make under behavioral and contextual targeting becomes identical. However, because we have the closed-form

probabilities, we can computationally examine the equilibrium in cases with more than two firms. Through extensive simulations with $K = 3, 4$ in the Online Appendix, we establish the same insight: There is no set of parameters under which the equilibrium under behavioral targeting is more polarizing than that under contextual targeting.

6.4. Endogenous Entry

In principle, privacy regulation might influence not just the content strategies of existing firms but also the number of firms operating in equilibrium. In this extension, we endogenize the number of firms, so that the equilibrium number of firms may change following a ban on behavioral ad targeting. To this end, we assume that there is a fixed entry cost $c > 0$ for each firm, and there are enough potential entrants such that new firms will always enter the market as long as they can make a nonnegative profit. When the entry cost c is too large, no firm operates in equilibrium, and the problem becomes uninteresting. Further, given that we have studied the monopoly and duopoly case in our main model, in this extension, we assume that c is small enough such that there is at least one firm at each position $x = 0, 1/2, 1$. In line with our baseline model, we maintain the assumption that each firm can only pick one position.[21]

Although our discrete modeling framework is sufficient for a small number of firms, it does not capture the complexities of a large market. The fundamental reason is that with three positions, firms cannot differentiate enough. This implies that the probability of a type θ consumer consuming content at position $x \neq \theta$ will exponentially decrease with more firms at a position, which in turn implies that contextual targeting will be identical to behavioral targeting, two properties at odds with the empirical reality of media markets. Therefore, we consider a position-specific rather than firm-specific idiosyncratic term in this section.[22] Lastly, we make a technical assumption on the number of firms in this extension: There is a continuum of firms rather than an integer number of firms. This assumption simplifies the analyses and has been adopted in seminal papers studying market entry (Hopenhayn 1992, Melitz 2003).[23]

In equilibrium, free entry and market clearing imply that each firm makes zero profits. Following privacy regulation, each consumer becomes less valuable to advertisers because uncertainty about the consumer's type leads to worse ad-matching quality, resulting in a lower total number of firms. Therefore, we will compare equilibrium outcomes with and without privacy regulation by the share of niche firms in each. We have the following result.

Proposition 9. *The consumer's choice of niche content is privacy-reducing over the choice of mainstream content. The ratio of the number of niche firms to the number of*

mainstream firms is higher under contextual ad targeting than under behavioral ad targeting.

As the proposition shows, the main insight in the main model extends to the case with endogenous entry when there are more firms and when the equilibrium number of firms depends endogenously on the privacy policy. Put differently, we find that with endogenous entry and privacy regulation, mainstream firms represent a smaller slice of a smaller pie. As a result, the decrease in the provision of centrist news is strengthened compared with the exogenous case.

7. Conclusion

The increased use of behavioral ad targeting by websites has heightened concerns by consumer privacy advocates and regulators. In this paper, we study the consequences of privacy regulations in the media landscape and ask the following question: How does a privacy regulation affect the equilibrium strategies of media firms in terms of ideological positioning and content polarization? We theoretically examine this question by building a model of product positioning and consumer demand. In our model, media firms choose their ideological positioning and create content, and consumers select their preferred content based on their ideology and idiosyncratic shocks. Depending on the eyeballs media firms generate, these firms extract advertising revenue based on the expected efficacy of their ads.

To examine the impact of privacy regulations, we focus on two primary information environments: behavioral and contextual ad targeting. Under *behavioral ad targeting*, media firms can perfectly track consumers and target their ads, reflecting the case without privacy regulation. In contrast, under *contextual ad targeting*, media firms cannot track readers and are therefore limited to tailoring ads to content, relying on the information contained in consumers' self-selection into the media firm's content. We show that banning behavioral ad targeting incentivizes media firms to shift toward niche (or polarizing) content strategies to improve their ability to draw inferences about consumers. In monopoly, this holds even when the increase in content polarization decreases total demand and consumer utility from content consumption. In duopoly, we document a stronger shift toward content polarization due to privacy regulation, although this shift does not reduce consumer utility from content consumption.

The findings of this study shed light on the complex connection between privacy and polarization. In recent years, many in the press have underscored the pivotal role played by increasingly precise personalization and behavioral ad targeting in the surge of political polarization within the United States. This perspective posits that hyper-personalization due to the lack of privacy regulation led to the creation of echo chambers and increased polarization. However, our research presents an alternative reversed link between privacy and content polarization. In our model, the inability to personalize content compels firms to use their content strategy as a targeting strategy that relies on consumers' self-selection into content. This strategic shift creates a unique incentive for producing more partisan content in both monopoly and duopoly scenarios.

Our paper contributes to the literature in several ways. First, we identify a counter-intuitive and theoretically robust link between privacy and polarization. We show that a privacy regulation leads to greater content polarization in equilibrium when media firms monetize based on advertising, even when content polarization reduces total demand and consumer utility from content consumption. Therefore, our paper identifies the firm's inference motive as another driver of content polarization beyond consumer demand and highlights the unintended consequences of privacy regulations in media markets. Second, we bring a novel inference-based framework to study the equilibrium effects of privacy regulation. In particular, our framework suggests that by limiting what firms *know* about consumers, privacy regulations shift firms' focus to what they *want to know* about consumers, thereby creating an incentive to use content strategies as means for more accurate inference. Put differently, we stress an important shift from viewing content polarization as a mere reflection of consumers' preexisting preferences, which is the case with personalization. Under contextual ad targeting, content consumed can become *more extreme and polarized* than consumers' preexisting preferences. Lastly, our paper offers important implications for managers and policymakers. In particular, we present the equilibrium effect of privacy regulations in media markets as a function of interpretable model parameters. We show that privacy regulations can have an unintended consequence of increased content polarization. Further, we propose that investments that increase content quality can mitigate the negative effects of privacy regulations on content polarization.

Nevertheless, our paper has certain limitations. First, we focused on discrete consumer types for tractability and interpretability reasons. Future research can extend our framework to a continuous setting. Second, in order to highlight the media firm's problem, we did not explicitly model the advertising marketplace. One could endogenize the advertising marketplace and obtain new insights for the advertising side of the market. Finally, our paper abstracted away from the dynamics between content polarization and opinion polarization. Future research could incorporate these dynamics and examine their long-run equilibrium.

Appendix

Proof of Lemma 1. Let p_j denote the posterior probability of $\theta = j$ conditional on the information $\mathcal{I}$, that is, $p_j = P(\theta = j | \mathcal{I})$. Now, we write the optimal ad choice as follows:

$$
\begin{aligned}
&\max_a \mathbf{E}[\mathcal{M}(a; \theta) | \mathcal{I}] \\
&= \max_a [1 - \gamma \mathbf{E}[(a - \theta)^2 | \mathcal{I}]] \\
&= \max_a \left[1 - \gamma \left(p_0 a^2 + p_{1/2} \left(a - \frac{1}{2} \right)^2 + p_1 (a - 1)^2 \right) \right].
\end{aligned}
$$

The first-order condition is $-\gamma \left(2 p_0 a + 2 p_{1/2} \left(a - \frac{1}{2} \right) + 2 p_1 (a - 1) \right)$ $= 0$, which results in $a^* = \frac{p_{1/2} + 2 p_1}{2} = \mathbf{E}[\theta | \mathcal{I}]$. $\square$

Proof of Corollary 1. The media firm's profit maximization problem is the following:

$$
\begin{aligned}
\max_{x, a} D(x) \cdot \mathbf{E}[\mathcal{M}(a; \theta) | \mathcal{I}] &= \max_x D(x) \cdot \mathbf{E}[\mathcal{M}(a^*; \theta) | \mathcal{I}] \\
&= \max_x D(x) \cdot \mathbf{E}[1 - \gamma (a^* - \theta)^2 | \mathcal{I}] \\
&= \max_x D(x) \cdot (1 - \gamma \mathbf{E}[(\mathbf{E}[\theta | \mathcal{I}] - \theta)^2 | \mathcal{I}]) \\
&= \max_x D(x) \cdot [1 - \gamma \mathbf{Var}(\theta | \mathcal{I})]. \quad \square
\end{aligned}
$$

Proof of Lemma 2. One can see that consumers whose type matches exactly the media firm's location will consume for sure. Therefore, we just need to determine whether other consumers prefer consuming the content to the outside option.

1. Mainstream content strategy $x = 1/2$: The demand from mainstream consumers is $1 - 2\lambda$. Consider type 0 consumers. Their utility from consuming the content is $U(1/2; 0) = \epsilon$. They will consume the content if and only if their consumption utility is greater than that of the outside option, i.e., $\epsilon > 0$. Therefore, the consumption probability is $P(\epsilon > 0) = 1 - F(0)$. Hence, the demand from type 0 consumers is $\lambda[1 - F(0)]$. By symmetry, the demand from type 1 consumers is also $\lambda[1 - F(0)]$. In sum, the total demand is $1 - 2\lambda + 2\lambda[1 - F(0)] = 1 - 2F(0)\lambda$.

2. Niche content strategy $x = 0$: The demand from type 0 consumers is λ and the demand from type 1 consumers is 0. Consider type 1/2 consumers. Their utility from consuming the content is $U(0; 1/2) = \epsilon$. They will consume the content if and only if $\epsilon > 0$, their utility from the outside option. By the same argument in the previous case, one can see that the demand from type 1/2 consumers is $(1 - 2\lambda)[1 - F(0)]$. In

sum, the total demand is $\lambda + (1 - 2\lambda)[1 - F(0)] = 1 - F(0) + [2F(0) - 1]\lambda$.

3. Niche content strategy $x = 1$: It is symmetric to the previous $x = 0$ case. $\square$

Proof of Proposition 1. Under behavioral ad targeting, the match between the ad and the consumer type is always 1. Therefore, the media firm chooses the content strategy that leads to the highest demand. According to Lemma 2, a mainstream positioning generates a higher demand than niche content does if and only if $1 - 2F(0)\lambda > 1 - F(0) + [2F(0) - 1]\lambda \Longleftrightarrow [4F(0) - 1]\lambda > F(0)$. $\square$

Proof of Lemma 3. Because of the symmetry between the two types of niche content, we only need to compare $\text{Privacy}(\{0\})$ and $\text{Privacy}(\{1/2\})$.

1. Mainstream content $x = 1/2$. According to Lemma 2, the demand from type 0 and from type 1 consumers are identical. Therefore, $\mathbf{E}[\theta | x = 1/2] = 1/2$.

$$
\text{Privacy}(\{1/2\}) = \mathbf{E}[(\theta - 1/2)^2 | x = 1/2] = \frac{1}{2} \cdot \frac{\lambda[1 - F(0)]}{1 - 2\lambda F(0)}
$$

2. Niche content $x = 0$.

$$
\begin{aligned}
\text{Privacy}(\{0\}) &= \mathbf{E}[\theta^2 | x = 0] - \mathbf{E}[\theta | x = 0]^2 \\
&= \frac{[1 - F(0)](1 - 2\lambda)}{\lambda + [1 - F(0)](1 - 2\lambda)} \cdot \left(\frac{1}{2} \right)^2 \\
&\quad - \left\{ \frac{[1 - F(0)](1 - 2\lambda)}{1 - F(0) + [2F(0) - 1]\lambda} \cdot \frac{1}{2} \right\}^2 \\
&= \frac{1}{4} \cdot \frac{[1 - F(0)](1 - 2\lambda)\lambda}{\{\lambda + [1 - F(0)](1 - 2\lambda)\}^2}
\end{aligned}
$$

Therefore, $\text{Privacy}(\{0\}) < \text{Privacy}(\{1/2\})$

$$
\begin{aligned}
&\Longleftrightarrow \frac{1}{4} \cdot \frac{[1 - F(0)](1 - 2\lambda)\lambda}{\{\lambda + [1 - F(0)](1 - 2\lambda)\}^2} < \frac{1}{2} \cdot \frac{\lambda[1 - F(0)]}{1 - 2\lambda F(0)} \\
&\Longleftrightarrow \lambda + \lambda(1 - 2\lambda)F(0) + 2\lambda(1 - 2\lambda)[1 - F(0)] \\
&\quad + \lambda^2 + [1 - F(0)]^2 (1 - 2\lambda)^2 > \frac{1}{2}
\end{aligned}
$$

The LHS is greater than $\lambda + \lambda(1 - 2\lambda)F(0) + \lambda(1 - 2\lambda)[1 - F(0)] + \lambda^2 = \lambda(2 - \lambda)$. Because $\lambda \in (1/3, 1/2)$ and $\lambda(2 - \lambda)$ increases in λ for $\lambda < 1$, we have the LHS $> \lambda(2 - \lambda) > 1/3(2 - 1/3) = 5/9 > 1/2$. Therefore, $\text{Privacy}(\{0\}) < \text{Privacy}(\{1/2\})$ always holds. By definition, niche content is privacy-reducing over mainstream content. $\square$

Proof of Proposition 2. Consider contextual ad targeting. By symmetry, we only need to compare the firm's expected profits from type 0 niche content strategy and mainstream content strategy:

$$
\begin{aligned}
\pi^c(0) > \pi^c(1/2) &\Longleftrightarrow D(0) \cdot [1 - \gamma \mathbf{Var}(\theta | x = 0)] \\
&\quad > D(1/2) \cdot [1 - \gamma \mathbf{Var}(\theta | x = 1/2)] \\
&\Longleftrightarrow \gamma > \frac{D(1/2)/D(0) - 1}{\mathbf{Var}(\theta | x = 1/2)D(1/2)/D(0)} =: \gamma^m. \\
&\quad\quad\quad - \mathbf{Var}(\theta | x = 0)
\end{aligned}
$$

Because $D(1/2) > D(0)$ and $\mathbf{Var}[\theta | x = 0] < \mathbf{Var}[\theta | x = 1/2]$, we have $\gamma^m > 0$. Therefore, the media firm chooses niche content strategy if $\gamma > \gamma^m$ and mainstream content strategy if $\gamma < \gamma^m$ under contextual ad targeting. Lastly, we show that γ^m, which can be expressed as $4\{F(0) + [1 - 4F(0)]\lambda\}\{1 -$

$F(0) + [2F(0) - 1]\lambda\}/\{[1 - F(0)]\lambda[1 - 2F(0) + 4F(0)\lambda]\}$, decreases in λ.

1. If $F(0) < 1/4$, then $4\{1 - F(0) + [2F(0) - 1]\lambda\}$ in the numerator decreases in λ, whereas $[1 - F(0)][1 - 2F(0) + 4F(0)\lambda]$ in the denominator increases in λ. One can also see that the remaining part of γ^m, $\{F(0) + [1 - 4F(0)]\lambda\}/\lambda = F(0)/\lambda + 1 - 4F(0)$, decreases in λ.

2. If $1/4 \le F(0) \le 1/2$, then both $F(0) + [1 - 4F(0)]\lambda$ and $1 - F(0) + [2F(0) - 1]\lambda$ decreases in λ. Therefore, the numerator of γ^m decreases in λ. One can see that the denominator of γ^m increases in λ.

3. If $F(0) > 1/2$, then $4\{F(0) + [1 - 4F(0)]\lambda\}$ in the numerator decreases in λ, whereas $[1 - F(0)][1 - 2F(0) + 4F(0)\lambda]$ in the denominator increases in λ. One can also see that the remaining part of γ^m, $\{1 - F(0) + [2F(0) - 1]\lambda\}/\lambda = [1 - F(0)]/\lambda + 2F(0) - 1$, decreases in λ.   □

Proof of Proposition 3. We first compute the consumer welfare for each content strategy.

1. Mainstream content strategy $x = 1/2$: The consumer welfare is $(1 - 2\lambda)[1/2 + \mathbf{E}[\epsilon]] + 2[1 - F(0)]\lambda\mathbf{E}[\epsilon|\epsilon > 0] = 1/2 - \lambda + 2[1 - F(0)]\lambda\mathbf{E}[\epsilon|\epsilon > 0]$.

2. Niche content strategy $x = 0$: The consumer welfare is $\lambda[1/2 + \mathbf{E}[\epsilon]] + [1 - F(0)](1 - 2\lambda)\mathbf{E}[\epsilon|\epsilon > 0] = \lambda/2 + [1 - F(0)](1 - 2\lambda)\mathbf{E}[\epsilon|\epsilon > 0]$.

3. Niche content strategy $x = 1$: It is symmetric to the previous $x = 0$ case.

Niche content strategy leads to lower consumer welfare if and only if

$$\lambda/2 + [1 - F(0)](1 - 2\lambda)\mathbf{E}[\epsilon|\epsilon > 0]$$
$$< 1/2 - \lambda + 2[1 - F(0)]\lambda\mathbf{E}[\epsilon|\epsilon > 0]$$
$$\Longleftrightarrow \lambda < \hat{\lambda} = \frac{1 - 2[1 - F(0)]\mathbf{E}[\epsilon|\epsilon > 0]}{3 - 8[1 - F(0)]\mathbf{E}[\epsilon|\epsilon > 0]}.$$

In addition, $\hat{\lambda} > 1/3 \Longleftrightarrow 2[1 - F(0)]\mathbf{E}[\epsilon|\epsilon > 0] > 0$, which holds for any distribution of ϵ.   □

Proof of Proposition 4. We first simplify the demand in Lemma 4 when $\epsilon \sim U[-1/2, 1/2]$.

Corollary A.1. *Suppose $\epsilon \sim U[-1/2, 1/2]$. The demand under each possible equilibrium is as follows:*

1. *Under (0,1) strategy profile, Firm 1's total demand is $3/8 + \lambda/4$. Among them, λ are type 0 consumers and $3/8 - 3\lambda/4$ are mainstream consumers. Firm 2's total demand is $3/8 + \lambda/4$. Among them, λ are type 1 consumers and $3/8 - 3\lambda/4$ are mainstream consumers.*

2. *Under (1/2,1/2) strategy profile, each firm's total demand is $1/2 - \lambda/4$. Among them, $1/2 - \lambda$ are mainstream consumers, $3\lambda/8$ are type 0 consumers, and $3\lambda/8$ are type 1 consumers.*

3. *Under (0,1/2) strategy profile, Firm 1's total demand is $1/8 + 5\lambda/8$. Among them, $7\lambda/8$ are type 0 consumers and $(1 - 2\lambda)/8$ are mainstream consumers. Firm 2's total demand is $7/8 - 9\lambda/8$. Among them, $\lambda/8$ are type 0 consumers, $\frac{7}{8}(1 - 2\lambda)$ are type 1/2 consumers, and $\lambda/2$ are type 1 consumers.*

Now consider the equilibrium strategy under behavioral ad targeting. Under behavioral ad targeting, the advertising revenue per consumer is one. Therefore, each firm's profit equals its total demand. Suppose the equilibrium is (0,1/2). Firm 1 will deviate from $x = 0$ to $x = 1/2$ if and only if

$$1/2 - \lambda/4 > 1/8 + 5\lambda/8 \Longleftrightarrow \lambda < \lambda_2 = 3/7. \quad (A.1)$$

Firm 2 will deviate from $x = 1/2$ to $x = 1$ if and only if

$$3/8 + \lambda/4 > 7/8 - 9\lambda/8 \Longleftrightarrow \lambda > \lambda_1 = 4/11. \quad (A.2)$$

From the above conditions and noting that $\lambda_1 < \lambda_2$, one can see that (0,1/2) will never be an equilibrium.

Now suppose the equilibrium is (1/2,1/2). Equation (A.1) implies that the firm will deviate to zero or one if $\lambda > \lambda_2$. Similarly, if the equilibrium is (0,1), Equation (A.2) implies that the firm will deviate to 1/2 if $\lambda < \lambda_1$.

In sum, the equilibrium is (1/2,1/2) if $\lambda < \lambda_1$, (1/2,1/2) or (0,1) if $\lambda_1 < \lambda < \lambda_2$, and (0,1) if $\lambda > \lambda_2$.   □

Proof of Proposition 5. We first characterize the equilibrium ad choices and profits for each firm under contextual ad targeting.

1. Under (1/2,1/2) strategy profile, by symmetry, $a = \mathbf{E}(\theta) = 1/2$. $\mathbf{Var}(\theta) = \mathbf{E}[\theta - \mathbf{E}(\theta)]^2 = 3\lambda/[4(2 - \lambda)]$.

$$\pi(1/2, 1/2) = D(1/2, 1/2) \cdot [1 - \gamma\mathbf{Var}(\theta)] = \frac{2 - \lambda}{4} - \frac{3\gamma\lambda}{16}$$

2. Under (0,1) strategy profile, by symmetry, we only need to consider firm 1. Firm 2's ad choice will be $1 - a_1$, and firm 2's profit will be identical to firm 1's. We have $a_1 = \mathbf{E}(\theta) = 3(1 - 2\lambda)/[2(3 + 2\lambda)]$, $\mathbf{Var}(\theta) = \mathbf{E}[\theta - \mathbf{E}(\theta)]^2 = 6\lambda(1 - 2\lambda)/(3 + 2\lambda)^2$.

$$\pi(0, 1) = D(0, 1) \cdot [1 - \gamma\mathbf{Var}(\theta)] = \frac{2\lambda + 3}{8} - \frac{3\gamma\lambda(1 - 2\lambda)}{4(2\lambda + 3)}$$

3. Under (0,1/2) strategy profile, consider firm 1 first. We have $a_1 = \mathbf{E}(\theta) = \mathbf{E}_1(\theta) = (1 - 2\lambda)/[2(1 + 5\lambda)]$, $\mathbf{Var}_1(\theta) = \mathbf{E}_1[\theta - \mathbf{E}_1(\theta)]^2 = 7\lambda(1 - 2\lambda)/[4(1 + 5\lambda)^2]$.

$$\pi_1(0, 1/2) = D_1(0, 1/2) \cdot [1 - \gamma\mathbf{Var}_1(\theta)] = \frac{1 + 5\lambda}{8} - \frac{7\gamma\lambda(1 - 2\lambda)}{32(1 + 5\lambda)}$$

Consider firm 2 then. We have $a_2 = \mathbf{E}_2(\theta) = (7 - 6\lambda)/[2(7 - 9\lambda)]$, $\mathbf{Var}_2(\theta) = \mathbf{E}_2[\theta - \mathbf{E}_2(\theta)]^2 = (35\lambda - 54\lambda^2)/[4(7 - 9\lambda)^2]$.

$$\pi_2(0, 1/2) = D_2(0, 1/2) \cdot [1 - \gamma\mathbf{Var}_2(\theta)]$$
$$= \frac{7 - 9\lambda}{8} - \frac{\gamma(35\lambda - 54\lambda^2)}{32(7 - 9\lambda)}$$

Suppose the equilibrium is (0,1/2). Firm 1 will deviate from $x = 0$ to $x = 1/2$ if and only if

$$\pi(1/2, 1/2) > \pi_1(0, 1/2) \Longleftrightarrow \gamma < \gamma^d = (\lambda_2 - \lambda)\frac{28(1 + 5\lambda)}{\lambda(44\lambda - 1)}. \quad (A.3)$$

An immediate implication is that this deviation never happens if $\lambda \ge \lambda_2$.

Firm 2 will deviate from $x = 1/2$ to $x = 1$ if and only if[24]

$$\pi(0, 1) > \pi_2(0, 1/2) \Longleftrightarrow \gamma > \gamma^{d'}$$
$$= (\lambda_1 - \lambda)\frac{44(2\lambda + 3)(7 - 9\lambda)}{\lambda(-540\lambda^2 + 460\lambda - 63)}. \quad (A.4)$$

An immediate implication is that deviation always happens if $\lambda \ge \lambda_1$.[25] Therefore, (0,1/2) may only be an equilibrium if $\lambda < \lambda_1$. Now suppose $\lambda < \lambda_1$. In this case, we have shown that firm 1 will deviate if $\gamma < \gamma^d$ and firm 2 will deviate if $\gamma > \gamma^{d'}$. Some calculation yields that there exists a unique $\lambda_0 \in (1/3, \lambda_1)$ such that $\gamma^d < (>)\gamma^{d'}$ if $\lambda < (>)\lambda_0$. Therefore, at least

Table A.1. Complete Equilibria Under Behavioral and Contextual Ad Targeting in Duopoly

	Equilibrium under	
Condition	Behavioral ad targeting	Contextual ad targeting
$1/3 < \lambda < \lambda_1$ & $\gamma < \min\{\gamma^d, \gamma^{d'}\}$	$(1/2, 1/2)$	$(1/2, 1/2)$
$1/3 < \lambda < \lambda_0$ & $\gamma^d < \gamma < \gamma^{d'}$	$(1/2, 1/2)$	$(0, 1/2)$
$\lambda_0 < \lambda < \lambda_1$ & $\gamma^{d'} < \gamma < \gamma^d$	$(1/2, 1/2)$	$(1/2, 1/2)$ or $(0,1)$
$1/3 < \lambda < \lambda_1$ & $\gamma > \max\{\gamma^d, \gamma^{d'}\}$	$(1/2, 1/2)$	$(0,1)$
$\lambda_0 < \lambda_1 < \lambda < \lambda_2$ & $\gamma < \gamma^d$	$(1/2, 1/2)$ or $(0,1)$	$(1/2, 1/2)$ or $(0,1)$
$\lambda_1 < \lambda < \lambda_2$ & $\gamma > \gamma^d$	$(1/2, 1/2)$ or $(0,1)$	$(0,1)$
$\lambda > \lambda_2$	$(0,1)$	$(0,1)$

one deviation happens and $(0, 1/2)$ will not be an equilibrium if $\lambda > \lambda_0$. There is no deviation and $(0, 1/2)$ is an equilibrium if $\lambda < \lambda_0$ and $\gamma^d < \gamma < \gamma^{d'}$.

Now suppose the equilibrium is $(1/2, 1/2)$. Equation (A.3) implies that the firm will deviate to zero or one if $\gamma > \gamma^d$, which may hold if $\lambda < \lambda_2$ and always holds if $\lambda > \lambda_2$. Similarly, if the equilibrium is $(0, 1)$, Equation (A.4) implies that the firm will deviate to $1/2$ if $\gamma < \gamma^{d'}$, which may hold if $\lambda < \lambda_1$ and never holds if $\lambda > \lambda_1$. Therefore, $(0, 1)$ is an equilibrium if $\lambda > \lambda_1$ or if $\lambda < \lambda_1$ & $\gamma > \gamma^{d'}$. $(1/2, 1/2)$ is an equilibrium if $\lambda < \lambda_2$ & $\gamma < \gamma^d$.

We summarize the equilibria in Table A.1. □

Endnotes

[1] We thank one of the anonymous reviewers for suggesting this reference.

[2] In a different setting, Iyer and Yoganarasimhan (2021) propose a separate mechanism, where polarization is driven by individuals' strategic considerations in a group decision-making problem.

[3] Our use of "niche" and "mainstream" is in line with Johnson and Myatt (2006) and needs not have a market share interpretation.

[4] It is worth noting that our qualitative insight holds for $\lambda < 1/3$. The reason we focus on $\lambda \in [1/3, 1/2]$ is that the insights from these cases have clearer implications for polarization.

[5] Even though the mainstream-type consumer is the minority, the demand for mainstream content can be more than that for niche content under both monopoly and duopoly. This is because the consumer faces some idiosyncratic shock in content consumption (the error term of the utility function in Equation (1) on the next page). On average, a consumer is farther from a niche position than a mainstream one. Therefore, mainstream content attracts more consumers of a different type than niche content does. As we will show in Section 5, the equilibrium content strategy may be $1/2$ in the monopoly case and $(1/2, 1/2)$ in the duopoly case.

[6] The intuitive reason behind this restriction is the reputational cost that the firm would face by holding multiple ideological positions at once. It is worth emphasizing that we do not focus on news aggregators as they are only distributors of content and not creators.

[7] The assumption that only one content unit is consumed is without loss of generality. As long as individual consumption does not depend on θ, all our findings remain unchanged.

[8] The inclusion of the idiosyncratic term is justified by two important empirical observations. First, there is substantial evidence that

consumers switch between media firms different in ideological positions (Gentzkow and Shapiro 2011). Second, without a stochastic component, the choice of media firm fully reveals the consumer type, making behavioral and contextual ad targeting the same, which runs counter to most empirical findings that behavioral ad targeting yield substantially higher revenues (Goldfarb and Tucker 2011, Johnson 2022).

[9] In the general base utility extension in Section 6.1, the consumer located at one extreme may consume content from the opposite extreme and the main insights are robust.

[10] We remark that this formulation essentially reduces consumer types to one dimensional or, in other words, assumes that ideological and advertising preferences are perfectly correlated and abstract away from ads that are not ideological. Although this is a simplification motivated by analytical convenience, our modeling framework maps well into settings where ad targeting is based on the corresponding ideology. The prime example of such contexts is political advertising (e.g., donations) where the ad value depends on the political ideology of consumers. However, the modeling framework extends to other cases where such tribal segments exist. To that end, a more polarized consumption pattern, as shown in Schoenmueller et al. (2023), creates more value in ideology-based targeting, although we acknowledge that the correlation, despite being positive, is not perfect in those cases.

[11] The prior literature suggests the possibility of thin markets under behavioral ad targeting (Levin and Milgrom 2010, Zhang and Katona 2012, Rafieian and Yoganarasimhan 2021). However, the assumption of a perfectly competitive ad market is reasonable in our context because we have three broad consumer types. As such, if the market is thick enough for each, our assumption holds. Further, even in the presence of a thin market, the media firm can still extract a higher revenue from a better ad match by setting a reservation price, which is fundamentally the assumption we make.

[12] Although we consider the case where ad allocation happens after consumer's content choice to be more consistent with the real practice of online ad allocation, it is important to emphasize that none of our results would change if we advertising decisions are made after the ideological positioning decision but before consumers' content choices, as long as these decisions are made for each possible scenario in $\mathcal{I}$.

[13] Our definition of privacy is similar to the existing measures of privacy, such as differential privacy (Dwork and Roth 2014), in capturing the uncertainty about an individual. However, differential privacy measures generally provide the definition over a database and "seeks to guarantee to the potential participant that, irrespective of her participation decision, almost the same things can be learned from the released outcome," thereby granting individuals plausible deniability and offering researchers a worst-case guarantee (Heffetz and Ligett 2014, p. 82). Because our work is not concerned with the identifiability of individual records from a database, we use a natural variance-based definition.

[14] In other words, with a different functional form for $\mathcal{M}(a, \theta)$, the firm might find it optimal to deviate from $a^* = \mathbf{E}[\theta|\mathcal{I}]$ in the contextual ad targeting case.

[15] The first derivative is $2[1 - F(0)]/(8[1 - F(0)]\mathbf{E}[\epsilon|\epsilon > 0) - 3)^2$, which is always nonnegative.

[16] We show in the proof of Proposition 3 that $\hat{\lambda} > 1/3$ for any distribution of the idiosyncratic term ϵ.

[17] The extent of media polarization increases from $(1/2, 1/2)$ to $(0, 1/2)$ to $(0, 1)$.

[18] Table A.1 presents the equilibria for all the cases.

[19] This complementarity motive arises as long as at least one firm chooses niche content. When both firms choose mainstream strategies,

(1/2, 1/2), each firm's demand is simply cut in half, and their inferential ability is the same as a monopoly, as no additional consumer ideological self-selection arises.

[20] Analogously, we can consider an alternative specification where the consumer gains more utility from better-matched ads, which does not change the results. We chose the current setting because there is more evidence that consumers dislike advertising rather than enjoying seeing an ad.

[21] If a firm offers multiple products at different positions, then each product can be viewed as a separate "firm" in this setting.

[22] Our qualitative insight that the equilibrium under contextual ad targeting is at least as polarizing as that under behavioral ad targeting still holds under the position-specific idiosyncratic term in the duopoly case, as in the main model.

[23] Including the integrality constraint in the model does not provide any relevant economic insights.

[24] Firm 2 has a stronger incentive to deviate to $x = 1$ rather than $x = 0$ to soften competition. Therefore, we only need to consider its deviation to $x = 1$.

[25] One can check that $44(2\lambda + 3)(7 - 9\lambda)/[\lambda(-540\lambda^2 + 460\lambda - 63)] > 0$ for $\lambda \in (1/3, 1/2)$.

References

Acquisti A, Varian HR (2005) Conditioning prices on purchase history. *Marketing Sci.* 24(3):367–381.

Alcobendas M, Kobayashi S, Shum M (2021) The impact of privacy measures on online advertising markets. Preprint, submitted February 18, https://doi.org/10.2139/ssrn.3782889.

Amaldoss W, Du J (2023) How can publishers collaborate and compete with news aggregators? *J. Marketing Res.* 60(6):1114–1134.

Amaldoss W, Du J, Shin W (2021) Media platforms' content provision strategies and sources of profits. *Marketing Sci.* 40(3):527–547.

Amaldoss W, Du J, Shin W (2024) Pricing strategy of competing media platforms. *Marketing Sci.* 43(3):488–505.

Amaldoss W, Jerath K, Sayedi A (2016) Keyword management costs and "broad match" in sponsored search advertising. *Marketing Sci.* 35(2):259–274.

Ambrus A, Calvano E, Reisinger M (2016) Either or both competition: A "two-sided" theory of advertising with overlapping viewerships. *Amer. Econom. J. Microeconom.* 8(3):189–222.

Anderson SP, Coate S (2005) Market provision of broadcasting: A welfare analysis. *Rev. Econom. Stud.* 72(4):947–972.

Athey S, Calvano E, Gans JS (2018) The impact of consumer multihoming on advertising markets and media competition. *Management Sci.* 64(4):1574–1590.

Bakshy E, Messing S, Adamic LA (2015) Exposure to ideologically diverse news and opinion on Facebook. *Science* 348(6239):1130–1132.

Bergemann D, Bonatti A (2011) Targeting in advertising markets: Implications for offline versus online media. *RAND J. Econom.* 42(3):417–443.

Bergemann D, Bonatti A (2015) Selling cookies. *Amer. Econom. J. Microeconom.* 7(3):259–294.

Bergemann D, Bonatti A, Smolin A (2018) The design and price of information. *Amer. Econom. Rev.* 108(1):1–48.

Berman R, Katona Z (2020) Curation algorithms and filter bubbles in social networks. *Marketing Sci.* 39(2):296–316.

Berman R, Oery A, Zheng X (2023) Influence or advertise: The role of social learning in influencer marketing. Preprint, submitted January 18, https://doi.org/10.2139/ssrn.4324888.

Berman R, Melumad S, Humphrey C, Meyer R (2019) A tale of two twitterspheres: Political microblogging during and after the 2016 primary and presidential debates. *J. Marketing Res.* 56(6):895–917.

Bonatti A, Huang Y, Villas-Boas J (2023) A theory of the effects of privacy. Technical report, Massachusetts Institute of Technology, Cambridge.

Boxell L, Gentzkow M, Shapiro JM (2017) Greater internet use is not associated with faster growth in political polarization among us demographic groups. *Proc. Natl. Acad. Sci. USA* 114(40):10612–10617.

Brynjolfsson E, Collis A, Liaqat A, Kutzman D, Garro H, Deisenroth D, Wernerfelt N (2024) The consumer welfare effects of online ads: Evidence from a 9-year experiment. NBER Working Paper No. 32846, National Bureau of Economic Research, Cambridge, MA.

Chaimanowong W, Cisternas F, Despotakis S (2023) The effects of microtargeting on conversion rates and the role of information asymmetry. Preprint, submitted August 2, https://doi.org/10.2139/ssrn.4525987.

Chen Y, Narasimhan C, Zhang ZJ (2001) Individual marketing with imperfect targetability. *Marketing Sci.* 20(1):23–41.

Choi WJ, Jerath K (2022) Privacy and consumer empowerment in online advertising. *Frontiers Marketing* 15(3):153–212.

Choi WJ, Sayedi A (2019) Learning in online advertising. *Marketing Sci.* 38(4):584–608.

Choi WJ, Sayedi A (2023) Open and private exchanges in display advertising. *Marketing Sci.* 42(3):451–475.

Choi WJ, Sayedi A (2024) Agency market power and information disclosure in online advertising. *Marketing Sci.* 43(6):1279–1298.

Choi WJ, Jerath K, Sarvary M (2020) Advertising and price competition under consumer data privacy choices. Preprint, submitted November 25, https://doi.org/10.2139/ssrn.3708273.

Choi WJ, Jerath K, Sarvary M (2023) Consumer privacy choices and (un) targeted advertising along the purchase journey. *J. Marketing Res.* 60(5):889–907.

Dukes A, Gal-Or E (2003) Negotiations and exclusivity contracts for advertising. *Marketing Sci.* 22(2):222–245.

Dwork C, Roth A (2014) The algorithmic foundations of differential privacy. *Frontiers Theoretical Comput. Sci.* 9(3–4):211–407.

Fainmesser IP, Galeotti A, Momot R (2023a) Consumer profiling via information design. Preprint, submitted December 20, http://dx.doi.org/10.2139/ssrn.4655468.

Fainmesser IP, Galeotti A, Momot R (2023b) Digital privacy. *Management Sci.* 69(6):3157–3173.

Flaxman S, Goel S, Rao JM (2016) Filter bubbles, echo chambers, and online news consumption. *Public Opinion Quart.* 80(S1):298–320.

Gal-Or E, Dukes A (2003) Minimum differentiation in commercial media markets. *Econom. Management Strategy* 12(3):291–325.

Gentzkow M, Shapiro JM (2006) Media bias and reputation. *J. Political Econom.* 114(2):280–316.

Gentzkow M, Shapiro JM (2010) What drives media slant? Evidence from US daily newspapers. *Econometrica* 78(1):35–71.

Gentzkow M, Shapiro JM (2011) Ideological segregation online and offline. *Quart. J. Econom.* 126(4):1799–1839.

Godes D, Ofek E, Sarvary M (2009) Content vs. advertising: The impact of competition on media firm strategy. *Marketing Sci.* 28(1):20–35.

Goldberg SG, Johnson GA, Shriver SK (2024) Regulating privacy online: An economic evaluation of the GDPR. *Amer. Econom. J. Econom. Policy* 16(1):325–358.

Goldfarb A, Tucker CE (2011) Privacy regulation and online advertising. *Management Sci.* 57(1):57–71.

Heffetz O, Ligett K (2014) Privacy and data-based research. *J. Econom. Perspectives* 28(2):75–98.

Hopenhayn HA (1992) Entry, exit, and firm dynamics in long run equilibrium. *Econometrica* 60(5):1127–1150.

Hosseinmardi H, Ghasemian A, Clauset A, Mobius M, Rothschild DM, Watts DJ (2021) Examining the consumption of radical content on YouTube. *Proc. Natl. Acad. Sci. USA* 118(32):1–8.

Hu M, Momot R, Wang J (2022) Privacy management in service systems. *Manufacturing Service Oper. Management* 24(5):2761–2779.

Iyer G, Singh S (2022) Persuasion contest: Disclosing own and rival information. *Marketing Sci.* 41(4):682–709.

Iyer G, Yoganarasimhan H (2021) Strategic polarization in group interactions. *J. Marketing Res.* 58(4):782–800.

Iyer G, Soberman D, Villas-Boas JM (2005) The targeting of advertising. *Marketing Sci.* 24(3):461–476.

Jain S, Qian K (2021) Compensating online content producers: A theoretical analysis. *Management Sci.* 67(11):7075–7090.

Johar M, Kumar N, Mookerjee V (2012) Content provision strategies in the presence of content piracy. *Inform. Systems Res.* 23(3-part-2):960–975.

Johnson G (2022) Economic research on privacy regulation: Lessons from the GDPR and beyond. Working paper, Boston University, Boston.

Johnson JP, Myatt DP (2006) On the simple economics of advertising, marketing, and product design. *Amer. Econom. Rev.* 96(3):756–784.

Johnson JP, Jungbauer T, Preuss M (2024) Online advertising, data sharing, and consumer control. *Management Sci.* 70(11):7984–8002.

Johnson GA, Shriver SK, Goldberg SG (2023) Privacy and market concentration: Intended and unintended consequences of the GDPR. *Management Sci.* 69(10):5695–5721.

Kamada Y, Kojima F (2014) Voter preferences, polarization, and electoral policies. *Amer. Econom. J. Microeconom.* 6(4):203–236.

Ke TT, Sudhir K (2023) Privacy rights and data security: GDPR and personal data markets. *Management Sci.* 69(8):4389–4412.

Ke TT, Shin J, Yu J (2023) A model of product portfolio design: Guiding consumer search through brand positioning. *Marketing Sci.* 42(6):1101–1124.

Kim E, Lelkes Y, McCrain J (2022) Measuring dynamic media bias. *Proc. Natl. Acad. Sci. USA* 119(32):e2202197119,

Kuksov D, Shachar R, Wang K (2013) Advertising and consumers' communications. *Marketing Sci.* 32(2):294–309.

Lauga DO, Ofek E, Katona Z (2022) When and how should firms differentiate? Quality and advertising decisions in a duopoly. *J. Marketing Res.* 59(6):1252–1265.

Laziuk E (2021) iOS 14.5 opt-in rate: Daily updates since launch. Accessed June 1, 2025, https://www.flurry.com/blog/ios-14-5-opt-in-rate-att-restricted-app-tracking-transparency-worldwide-us-daily-latest-update/.

Lei Y, Miao S, Momot R (2024) Privacy-preserving personalized revenue management. *Management Sci.* 70(7):4875–4892.

Levin J, Milgrom P (2010) Online advertising: Heterogeneity and conflation in market design. *Amer. Econom. Rev.* 100(2):603–607.

Lin S (2020) Two-sided price discrimination by media platforms. *Marketing Sci.* 39(2):317–338.

Mayzlin D, Shin J (2011) Uninformative advertising as an invitation to search. *Marketing Sci.* 30(4):666–685.

McClain C, Faverio M, Anderson M, Park E (2023) How Americans view data privacy. https://www.pewresearch.org/internet/2023/10/18/how-americans-view-data-privacy/.

Melitz MJ (2003) The impact of trade on intra-industry reallocations and aggregate industry productivity. *Econometrica* 71(6):1695–1725.

Miklós-Thal J, Goldfarb A, Haviv A, Tucker C (2024) Frontiers: Digital hermits. *Marketing Sci.* 43(4):697–708.

Momot R, Salikhov M (2022) Information elicitation from teams of privacy-conscious experts. Preprint, submitted October 14, https://doi.org/10.2139/ssrn.4248487.

Mullainathan S, Shleifer A (2005) The market for news. *Amer. Econom. Rev.* 95(4):1031–1053.

Ning ZE, Shin J, Yu J (2025) Targeted advertising as implicit recommendation: Strategic mistargeting and personal data opt-out. *Marketing Sci.* 44(2):390–410.

Novatiq (2024) The evolution of programmatic advertising: State of the nation. https://www.novatiq.com/the-evolution-of-programmatic-advertising-state-of-the-nation/.

Pariser E (2011) *The Filter Bubble: How the New Personalized Web Is Changing What We Read and How We Think* (Penguin, New York).

Rafieian O, Yoganarasimhan H (2021) Targeting and privacy in mobile advertising. *Marketing Sci.* 40(2):193–218.

Rhodes A, Zhou J (2024) Personalization and privacy choice. Preprint, submitted April 16, https://doi.org/10.2139/ssrn.4795118.

Ribeiro MH, Ottoni R, West R, Almeida VA, Meira W Jr (2020) Auditing radicalization pathways on YouTube. Hildebrandt M, Castillo C, eds. *Proc. Conf. Fairness Accountability Transparency* (Association for Computing Machinery, New York), 131–141.

Sayedi A (2018) Real-time bidding in online display advertising. *Marketing Sci.* 37(4):553–568.

Schoenmueller V, Netzer O, Stahl F (2023) Frontiers: Polarized America: From political polarization to preference polarization. *Marketing Sci.* 42(1):48–60.

Shin J, Shin W (2023) A theory of irrelevant advertising: An agency-induced targeting inefficiency. *Management Sci.* 69(8):4481–4497.

Shin J, Yu J (2021) Targeted advertising and consumer inference. *Marketing Sci.* 40(5):900–922.

Strack P, Yang KH (2024) Privacy-preserving signals. *Econometrica* 92(6):1907–1938.

Sun M, Zhu F (2013) Ad revenue and content commercialization: Evidence from blogs. *Management Sci.* 59(10):2314–2331.

Sweeney M (2023) How real-time bidding (RTB) changed online display advertising. https://clearcode.cc/blog/real-time-bidding-online-display-advertising/.

Taylor CR (2004) Consumer privacy and the market for customer information. *RAND J. Econom.* 35(4):631–650.

Tucker C, Zhang J (2010) Growing two-sided networks by advertising the user base: A field experiment. *Marketing Sci.* 29(5):805–814.

Villas-Boas JM (1999) Dynamic competition with customer recognition. *RAND J. Econom.* 30(4):604–631.

Villas-Boas JM (2004) Price cycles in markets with customer recognition. *RAND J. Econom.* 35(3):486–501.

Wilbur KC (2016) Advertising content and television advertising avoidance. *J. Media Econom.* 29(2):51–72.

Xiang Y, Sarvary M (2007) News consumption and media bias. *Marketing Sci.* 26(5):611–628.

Yang KH (2022) Selling consumer data for profit: Optimal market-segmentation design and its consequences. *Amer. Econom. Rev.* 112(4):1364–1393.

Yang J, Zhang J, Zhang Y (2025) Engagement that sells: Influencer video advertising on TikTok. *Marketing Sci.* 44(2):247–267.

Yao YJ (2024) Invitation to search or purchase? Optimal multi-attribute advertising. Preprint, submitted August 23, http://dx.doi.org/10.2139/ssrn.4936056.

Yao YJ (2025) Reputation for privacy. *Marketing Sci.*, ePub ahead of print April 16, https://doi.org/10.1287/mksc.2024.1006.

Zhang K, Katona Z (2012) Contextual advertising. *Marketing Sci.* 31(6):980–994.